

Homework 2

Zhangrui Huang (zh2680@nyu.edu)

Problem 1

(a). $L(w_1, w_2) = 0.5(aw_1^2 + bw_2^2)$, $\frac{\partial L}{\partial w_1} = aw_1$, $\frac{\partial L}{\partial w_2} = bw_2$, $\nabla L(w_1, w_2) = [aw_1, bw_2]$

To achieve the minimum value of L , set $\nabla L(w_1, w_2) = [aw_1, bw_2] = 0 \Rightarrow w_1 = 0, w_2 = 0$

(b). Gradient descent should be $w(t+1) = w(t) - \eta \frac{\partial L}{\partial w}$, where η is learning rate.

$$w_1(t+1) = w_1(t) - \eta \frac{\partial L}{\partial w_1} = w_1(t) - \eta a w_1(t) = (1 - \eta a) w_1(t) \Rightarrow \rho_1 = 1 - \eta a$$

$$w_2(t+1) = w_2(t) - \eta \frac{\partial L}{\partial w_2} = w_2(t) - \eta b w_2(t) = (1 - \eta b) w_2(t) \Rightarrow \rho_2 = 1 - \eta b$$

(c). To converge the result of gradient descent, $-1 < \rho_1, \rho_2 < 1$, so $0 < \eta < \min(\frac{2}{a}, \frac{2}{b})$

(b). a/b is a very large ratio, which means a is much larger than b . In this case even if η is optimally set for w_1 , it might be too small for w_2 , so the convergence rate of gradient descent for w_2 is very slow.

Problem 2

(a). The filter should be like

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

This is a vertical Sobel filter, it calculate the difference in intensity between the left and right pixels. Pixels in the middle are weighted more heavily. If there's a strong vertical edge, the difference in intensity will be high, and the output will be high.

(b). The filter should be like

$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}$$

This is a blurring filter, which is to replace each pixel value with the average value of its neighbors, including itself.

(c). The filter should be like

$$\begin{bmatrix} 0 & -1 & 0 \\ 0 & 5 & 0 \\ 0 & -1 & 0 \end{bmatrix}$$

This is a horizontal sharpening filter, it increases the weight of the central pixel while subtracting some of its horizontal neighbors' intensities. The result is that horizontal features get enhanced.

(d). The filter should be like

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

This is a noise reduction filter, it gives more weight to pixels closer to the center. It averages out the values in a pixel's neighborhood, which can help in reducing high-frequency noise.

Problem 3

(a). $IOU = \frac{A \cap B}{A \cup B}$ is a metric defined as the ratio of the area of intersection of two bounding boxes to the area of their union.

If there is no intersection between two boxes, the area of $A \cap B = 0, IOU = 0$

If box A equals to box B, $A \cap B = A \cup B, IOU = 1$

By definition, $A \cup B = A + B - A \cap B$, so $A \cup B \geq A \cap B, 0 \leq IOU \leq 1$

(b). When using the top-left and bottom-right coordinates to represent box,

$Box = (x_{top-left}, y_{top-left}, x_{bottom-right}, y_{bottom-right})$

The area of intersection, $A \cap B$, involves taking minimum and maximum operations to find the overlapping region.

$$x_{intersection-top-left} = \max(x_{top-left}^A, x_{top-left}^B)$$

$$y_{intersection-top-left} = \max(y_{top-left}^A, y_{top-left}^B)$$

$$x_{intersection-bottom-right} = \min(x_{bottom-right}^A, x_{bottom-right}^B)$$

$$y_{intersection-bottom-right} = \min(y_{bottom-right}^A, y_{bottom-right}^B)$$

Both the max and min operations are non-differentiable so IOU metric is non-differentiable and cannot be used as a loss function.

AlexNet

In this problem, you are asked to train a deep convolutional neural network to perform image classification. In fact, this is a slight variation of a network called *AlexNet*. This is a landmark model in deep learning, and arguably kickstarted the current (and ongoing, and massive) wave of innovation in modern AI when its results were first presented in 2012. AlexNet was the first real-world demonstration of a *deep* classifier that was trained end-to-end on data and that outperformed all other ML models thus far.

We will train AlexNet using the [CIFAR10](#) dataset, which consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. The classes are: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck.

A lot of the code you will need is already provided in this notebook; all you need to do is to fill in the missing pieces, and interpret your results.

Warning : AlexNet takes a good amount of time to train (~1 minute per epoch on Google Colab). So please budget enough time to do this homework.

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.optim.lr_scheduler import _LRScheduler
import torch.utils.data as data

import torchvision.transforms as transforms
import torchvision.datasets as datasets

from sklearn import decomposition
from sklearn import manifold
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

import copy
import random
import time
```

```
SEED = 1234

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True
```

Loading and Preparing the Data

Our dataset is made up of color images but three color channels (red, green and blue), compared to MNIST's black and white images with a single color channel. To normalize our data we need to calculate the means and standard deviations for each of the color channels independently, and normalize them.

```
ROOT = '.data'
train_data = datasets.CIFAR10(root = ROOT,
                             train = True,
                             download = True)
```

Files already downloaded and verified

```
# Compute means and standard deviations along the R,G,B channel

means = train_data.data.mean(axis = (0,1,2)) / 255
stds = train_data.data.std(axis = (0,1,2)) / 255
```

Next, we will do data augmentation. For each training image we will randomly rotate it (by up to 5 degrees), flip/mirror with probability 0.5, shift by +/-1 pixel. Finally we will normalize each color channel using the means/stds we calculated above.

```
train_transforms = transforms.Compose([
    transforms.RandomRotation(5),
    transforms.RandomHorizontalFlip(0.5),
    transforms.RandomCrop(32, padding = 2),
    transforms.ToTensor(),
    transforms.Normalize(mean = means,
                        std = stds)
])

test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean = means,
                        std = stds)
])
```

Next, we'll load the dataset along with the transforms defined above.

We will also create a validation set with 10% of the training samples. The validation set will be used to monitor loss along different epochs, and we will pick the model along the optimization path that performed the best, and report final test accuracy numbers using this model.

```
train_data = datasets.CIFAR10(ROOT,
                             train = True,
                             download = True,
                             transform = train_transforms)

test_data = datasets.CIFAR10(ROOT,
                            train = False,
                            download = True,
                            transform = test_transforms)
```

```
Files already downloaded and verified
Files already downloaded and verified
```

```

VALID_RATIO = 0.9

n_train_examples = int(len(train_data) * VALID_RATIO)
n_valid_examples = len(train_data) - n_train_examples

train_data, valid_data = data.random_split(train_data,
                                            [n_train_examples, n_valid_examples])

```

```

valid_data = copy.deepcopy(valid_data)
valid_data.dataset.transform = test_transforms

```

Now, we'll create a function to plot some of the images in our dataset to see what they actually look like.

Note that by default PyTorch handles images that are arranged `[channel, height, width]`, but `matplotlib` expects images to be `[height, width, channel]`, hence we need to permute the dimensions of our images before plotting them.

```

def plot_images(images, labels, classes, normalize = False):

    n_images = len(images)

    rows = int(np.sqrt(n_images))
    cols = int(np.sqrt(n_images))

    fig = plt.figure(figsize = (10, 10))

    for i in range(rows*cols):

        ax = fig.add_subplot(rows, cols, i+1)

        image = images[i]

        if normalize:
            image_min = image.min()
            image_max = image.max()
            image.clamp_(min = image_min, max = image_max)
            image.add_(-image_min).div_(image_max - image_min + 1e-5)

        ax.imshow(image.permute(1, 2, 0).cpu().numpy())
        ax.set_title(classes[labels[i]])
        ax.axis('off')

```

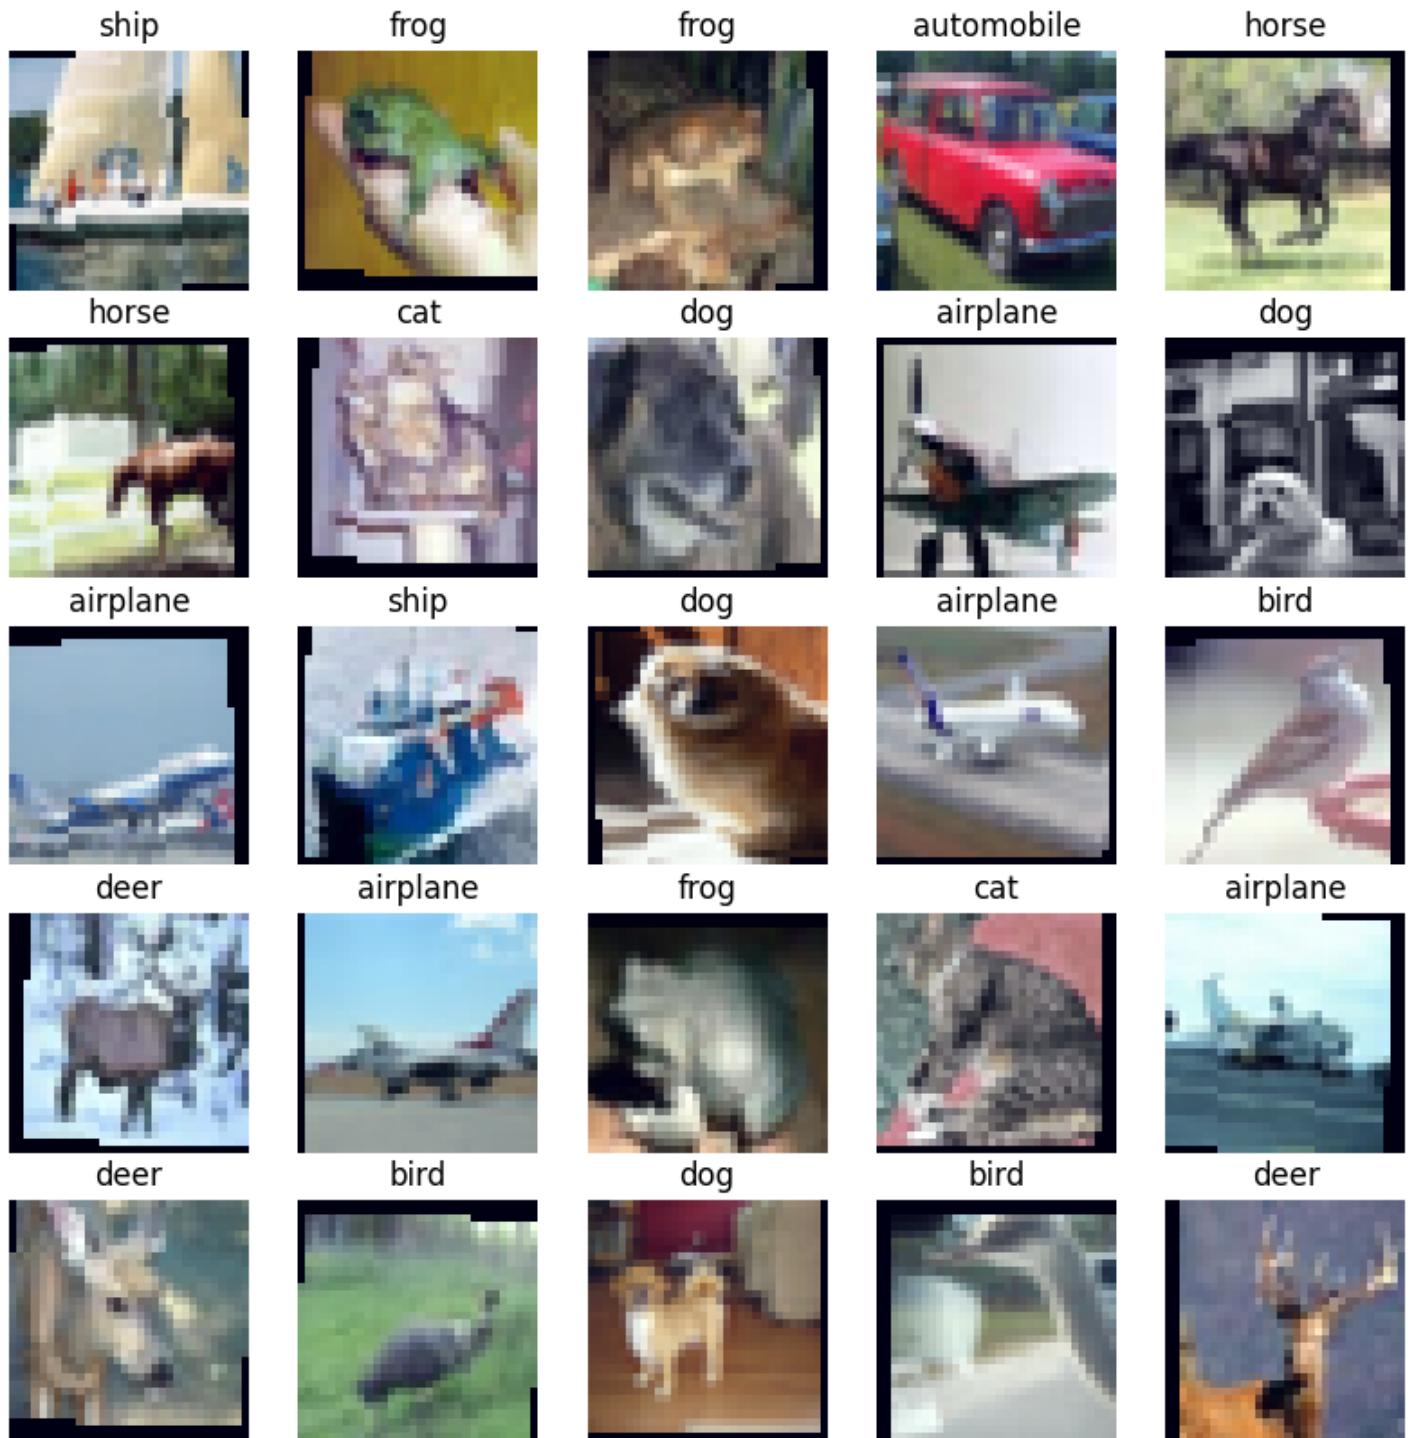
One point here: `matplotlib` is expecting the values of every pixel to be between $[0, 1]$, however our normalization will cause them to be outside this range. By default `matplotlib` will then clip these values into the $[0, 1]$ range. This clipping causes all of the images to look a bit weird - all of the colors are oversaturated. The solution is to normalize each image between $[0, 1]$.

```
N_IMAGES = 25

images, labels = zip(*[(image, label) for image, label in
                      [train_data[i] for i in range(N_IMAGES)]])

classes = test_data.classes
```

```
plot_images(images, labels, classes, normalize = True)
```



We'll be normalizing our images by default from now on, so we'll write a function that does it for us which we can use whenever we need to renormalize an image.

```

def normalize_image(image):
    image_min = image.min()
    image_max = image.max()
    image.clamp_(min = image_min, max = image_max)
    image.add_(-image_min).div_(image_max - image_min + 1e-5)
    return image

```

The final bit of the data processing is creating the iterators. We will use a large. Generally, a larger batch size means that our model trains faster but is a bit more susceptible to overfitting.

```

# Q1: Create data loaders for train_data, valid_data, test_data
# Use batch size 256

BATCH_SIZE = 256

train_iterator = data.DataLoader(train_data, shuffle=True, batch_size=BATCH_SIZE)

valid_iterator = data.DataLoader(valid_data, shuffle=False, batch_size=BATCH_SIZE)

test_iterator = data.DataLoader(test_data, shuffle=False, batch_size=BATCH_SIZE)

```

Defining the Model

Next up is defining the model.

AlexNet will have the following architecture:

- There are 5 2D convolutional layers (which serve as *feature extractors*), followed by 3 linear layers (which serve as the *classifier*).
- All layers (except the last one) have `ReLU` activations. (Use `inplace=True` while defining your ReLUs.)
- All convolutional filter sizes have kernel size 3×3 and padding 1.
- Convolutional layer 1 has stride 2. All others have the default stride (1).
- Convolutional layers 1,2, and 5 are followed by a 2D maxpool of size 2.
- Linear layers 1 and 2 are preceded by Dropouts with Bernoulli parameter 0.5.
- For the convolutional layers, the number of channels is set as follows. We start with 3 channels and then proceed like this:
 - $3 \rightarrow 64 \rightarrow 192 \rightarrow 384 \rightarrow 256 \rightarrow 256$

In the end, if everything is correct you should get a feature map of size $2 \times 2 \times 256 = 1024$.

- For the linear layers, the feature sizes are as follows:

- $1024 \rightarrow 4096 \rightarrow 4096 \rightarrow 10$.

(The 10, of course, is because 10 is the number of classes in CIFAR-10).

```

class AlexNet(nn.Module):
    def __init__(self, output_dim):
        super().__init__()

        self.features = nn.Sequential(
            # Define according to the steps described above
            nn.Conv2d(3, 64, kernel_size=3, stride=2, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2),

            nn.Conv2d(64, 192, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2),

            nn.Conv2d(192, 384, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            nn.Conv2d(384, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2)
        )

        self.classifier = nn.Sequential(
            # define according to the steps described above
            nn.Dropout(0.5),
            nn.Linear(1024, 4096),
            nn.ReLU(inplace=True),

            nn.Dropout(0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),

            nn.Linear(4096, output_dim)
        )

    def forward(self, x):
        x = self.features(x)
        h = x.view(x.shape[0], -1)
        x = self.classifier(h)
        return x, h

```

We'll create an instance of our model with the desired amount of classes.

```
OUTPUT_DIM = 10
model = AlexNet(OUTPUT_DIM)
```

Training the Model

We first initialize parameters in PyTorch by creating a function that takes in a PyTorch module, checking what type of module it is, and then using the `nn.init` methods to actually initialize the parameters.

For convolutional layers we will initialize using the *Kaiming Normal* scheme, also known as *He Normal*. For the linear layers we initialize using the *Xavier Normal* scheme, also known as *Glorot Normal*. For both types of layer we initialize the bias terms to zeros.

```
def initialize_parameters(m):
    if isinstance(m, nn.Conv2d):
        nn.init.kaiming_normal_(m.weight.data, nonlinearity = 'relu')
        nn.init.constant_(m.bias.data, 0)
    elif isinstance(m, nn.Linear):
        nn.init.xavier_normal_(m.weight.data, gain = nn.init.calculate_gain('relu'))
        nn.init.constant_(m.bias.data, 0)
```

We apply the initialization by using the model's `apply` method. If your definitions above are correct you should get the printed output as below.

```
model.apply(initialize_parameters)
```

```
AlexNet(
(features): Sequential(
(0): Conv2d(3, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
(1): ReLU(inplace=True)
(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(3): Conv2d(64, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(4): ReLU(inplace=True)
(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(7): ReLU(inplace=True)
(8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(9): ReLU(inplace=True)
(10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace=True)
(12): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Dropout(p=0.5, inplace=False)
```

```

(v). Dropout(p=0.5, inplace=False),
(1): Linear(in_features=1024, out_features=4096, bias=True)
(2): ReLU(inplace=True)
(3): Dropout(p=0.5, inplace=False)
(4): Linear(in_features=4096, out_features=4096, bias=True)
(5): ReLU(inplace=True)
(6): Linear(in_features=4096, out_features=10, bias=True)
)
)

```

We then define the loss function we want to use, the device we'll use and place our model and criterion on to our device.

```

optimizer = optim.Adam(model.parameters(), lr = 1e-3)
device = torch.device('cuda' if torch.cuda.is_available() else 'mps')
criterion = nn.CrossEntropyLoss()

model = model.to(device)
criterion = criterion.to(device)

```

This is formatted as code

We define a function to calculate accuracy...

```

def calculate_accuracy(y_pred, y):
    top_pred = y_pred.argmax(1, keepdim = True)
    correct = top_pred.eq(y.view_as(top_pred)).sum()
    acc = correct.float() / y.shape[0]
    return acc

```

As we are using dropout we need to make sure to "turn it on" when training by using `model.train()`.

```

def train(model, iterator, optimizer, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.train()

```

```

for (x, y) in iterator:

    x = x.to(device)
    y = y.to(device)

    optimizer.zero_grad()

    y_pred, _ = model(x)

    loss = criterion(y_pred, y)

    acc = calculate_accuracy(y_pred, y)

    loss.backward()

    optimizer.step()

    epoch_loss += loss.item()
    epoch_acc += acc.item()

return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

We also define an evaluation loop, making sure to "turn off" dropout with `model.eval()`.

```

def evaluate(model, iterator, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():

        for (x, y) in iterator:

            x = x.to(device)
            y = y.to(device)

            y_pred, _ = model(x)

            loss = criterion(y_pred, y)

            acc = calculate_accuracy(y_pred, y)

            epoch_loss += loss.item()
            epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

Next, we define a function to tell us how long an epoch takes.

```
def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs
```

Then, finally, we train our model.

Train it for 25 epochs (using the train dataset). At the end of each epoch, compute the validation loss and keep track of the best model. You might find the command `torch.save` helpful.

At the end you should expect to see validation losses of ~76% accuracy.

```
# Q3: train your model here for 25 epochs.
# Print out training and validation loss/accuracy of the model after each epoch
# Keep track of the model that achieved best validation loss thus far.

EPOCHS = 25

# Fill training code here

best_valid_loss = float('inf')

for epoch in range(EPOCHS):

    start_time = time.time()

    # Training
    train_loss, train_acc = train(model, train_iterator, optimizer, criterion, device)

    # Validation
    valid_loss, valid_acc = evaluate(model, valid_iterator, criterion, device)

    end_time = time.time()

    epoch_mins, epoch_secs = epoch_time(start_time, end_time)

    # Saving the best model
    if valid_loss < best_valid_loss:
        best_valid_loss = valid_loss
        torch.save(model.state_dict(), 'best-model.pt')

    # Printing results after the current epoch
    print(f'Epoch: {epoch+1:02} | Epoch Time: {epoch_mins}m {epoch_secs}s')
    print(f'\tTrain Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}%')
```

```
print(f'\t Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')
```

Epoch: 01 | Epoch Time: 0m 17s
Train Loss: 2.412 | Train Acc: 21.50%
Val. Loss: 1.615 | Val. Acc: 37.16%
Epoch: 02 | Epoch Time: 0m 10s
Train Loss: 1.547 | Train Acc: 42.50%
Val. Loss: 1.385 | Val. Acc: 50.26%
Epoch: 03 | Epoch Time: 0m 10s
Train Loss: 1.373 | Train Acc: 50.10%
Val. Loss: 1.243 | Val. Acc: 54.74%
Epoch: 04 | Epoch Time: 0m 10s
Train Loss: 1.274 | Train Acc: 54.00%
Val. Loss: 1.160 | Val. Acc: 59.15%
Epoch: 05 | Epoch Time: 0m 10s
Train Loss: 1.183 | Train Acc: 57.64%
Val. Loss: 1.097 | Val. Acc: 60.51%
Epoch: 06 | Epoch Time: 0m 10s
Train Loss: 1.123 | Train Acc: 60.36%
Val. Loss: 1.059 | Val. Acc: 62.83%
Epoch: 07 | Epoch Time: 0m 10s
Train Loss: 1.068 | Train Acc: 62.47%
Val. Loss: 1.000 | Val. Acc: 64.66%
Epoch: 08 | Epoch Time: 0m 10s
Train Loss: 1.014 | Train Acc: 64.19%
Val. Loss: 0.960 | Val. Acc: 66.28%
Epoch: 09 | Epoch Time: 0m 10s
Train Loss: 0.973 | Train Acc: 65.82%
Val. Loss: 0.932 | Val. Acc: 67.35%
Epoch: 10 | Epoch Time: 0m 10s
Train Loss: 0.937 | Train Acc: 67.23%
Val. Loss: 0.903 | Val. Acc: 68.04%
Epoch: 11 | Epoch Time: 0m 10s
Train Loss: 0.913 | Train Acc: 68.11%
Val. Loss: 0.860 | Val. Acc: 70.45%
Epoch: 12 | Epoch Time: 0m 10s
Train Loss: 0.871 | Train Acc: 69.58%
Val. Loss: 0.894 | Val. Acc: 68.85%
Epoch: 13 | Epoch Time: 0m 10s
Train Loss: 0.854 | Train Acc: 70.49%
Val. Loss: 0.838 | Val. Acc: 70.89%
Epoch: 14 | Epoch Time: 0m 10s
Train Loss: 0.829 | Train Acc: 71.00%
Val. Loss: 0.799 | Val. Acc: 72.94%
Epoch: 15 | Epoch Time: 0m 10s
Train Loss: 0.801 | Train Acc: 72.30%
Val. Loss: 0.809 | Val. Acc: 71.90%
Epoch: 16 | Epoch Time: 0m 10s

```
Train Loss: 0.785 | Train Acc: 72.90%
Val. Loss: 0.776 | Val. Acc: 73.65%
Epoch: 17 | Epoch Time: 0m 10s
Train Loss: 0.763 | Train Acc: 73.80%
Val. Loss: 0.794 | Val. Acc: 72.64%
Epoch: 18 | Epoch Time: 0m 10s
Train Loss: 0.758 | Train Acc: 73.87%
Val. Loss: 0.789 | Val. Acc: 73.25%
Epoch: 19 | Epoch Time: 0m 10s
Train Loss: 0.735 | Train Acc: 74.63%
Val. Loss: 0.762 | Val. Acc: 74.65%
Epoch: 20 | Epoch Time: 0m 10s
Train Loss: 0.715 | Train Acc: 75.22%
Val. Loss: 0.731 | Val. Acc: 75.45%
Epoch: 21 | Epoch Time: 0m 10s
Train Loss: 0.705 | Train Acc: 75.87%
Val. Loss: 0.737 | Val. Acc: 74.67%
Epoch: 22 | Epoch Time: 0m 10s
Train Loss: 0.683 | Train Acc: 76.59%
Val. Loss: 0.731 | Val. Acc: 75.41%
Epoch: 23 | Epoch Time: 0m 10s
Train Loss: 0.672 | Train Acc: 76.89%
Val. Loss: 0.733 | Val. Acc: 75.56%
Epoch: 24 | Epoch Time: 0m 10s
Train Loss: 0.660 | Train Acc: 77.11%
Val. Loss: 0.728 | Val. Acc: 75.30%
Epoch: 25 | Epoch Time: 0m 10s
Train Loss: 0.657 | Train Acc: 77.21%
Val. Loss: 0.728 | Val. Acc: 75.60%
```

Evaluating the model

We then load the parameters of our model that achieved the best validation loss. You should expect to see ~75% accuracy of this model on the test dataset.

Finally, plot the confusion matrix of this model and comment on any interesting patterns you can observe there. For example, which two classes are confused the most?

```
# Q4: Load the best performing model, evaluate it on the test dataset, and print test
accuracy.

# Also, print out the confusion matrox.
```

```
def get_predictions(model, iterator, device):
```

```
model.eval()

labels = []
probs = []

# Q4: Fill code here.
with torch.no_grad():

    for (x, y) in iterator:

        x = x.to(device)
        y = y.to(device)

        y_pred, _ = model(x)

        labels.append(y)
        probs.append(y_pred)

labels = torch.cat(labels, dim = 0)
probs = torch.cat(probs, dim = 0)

return labels, probs
```

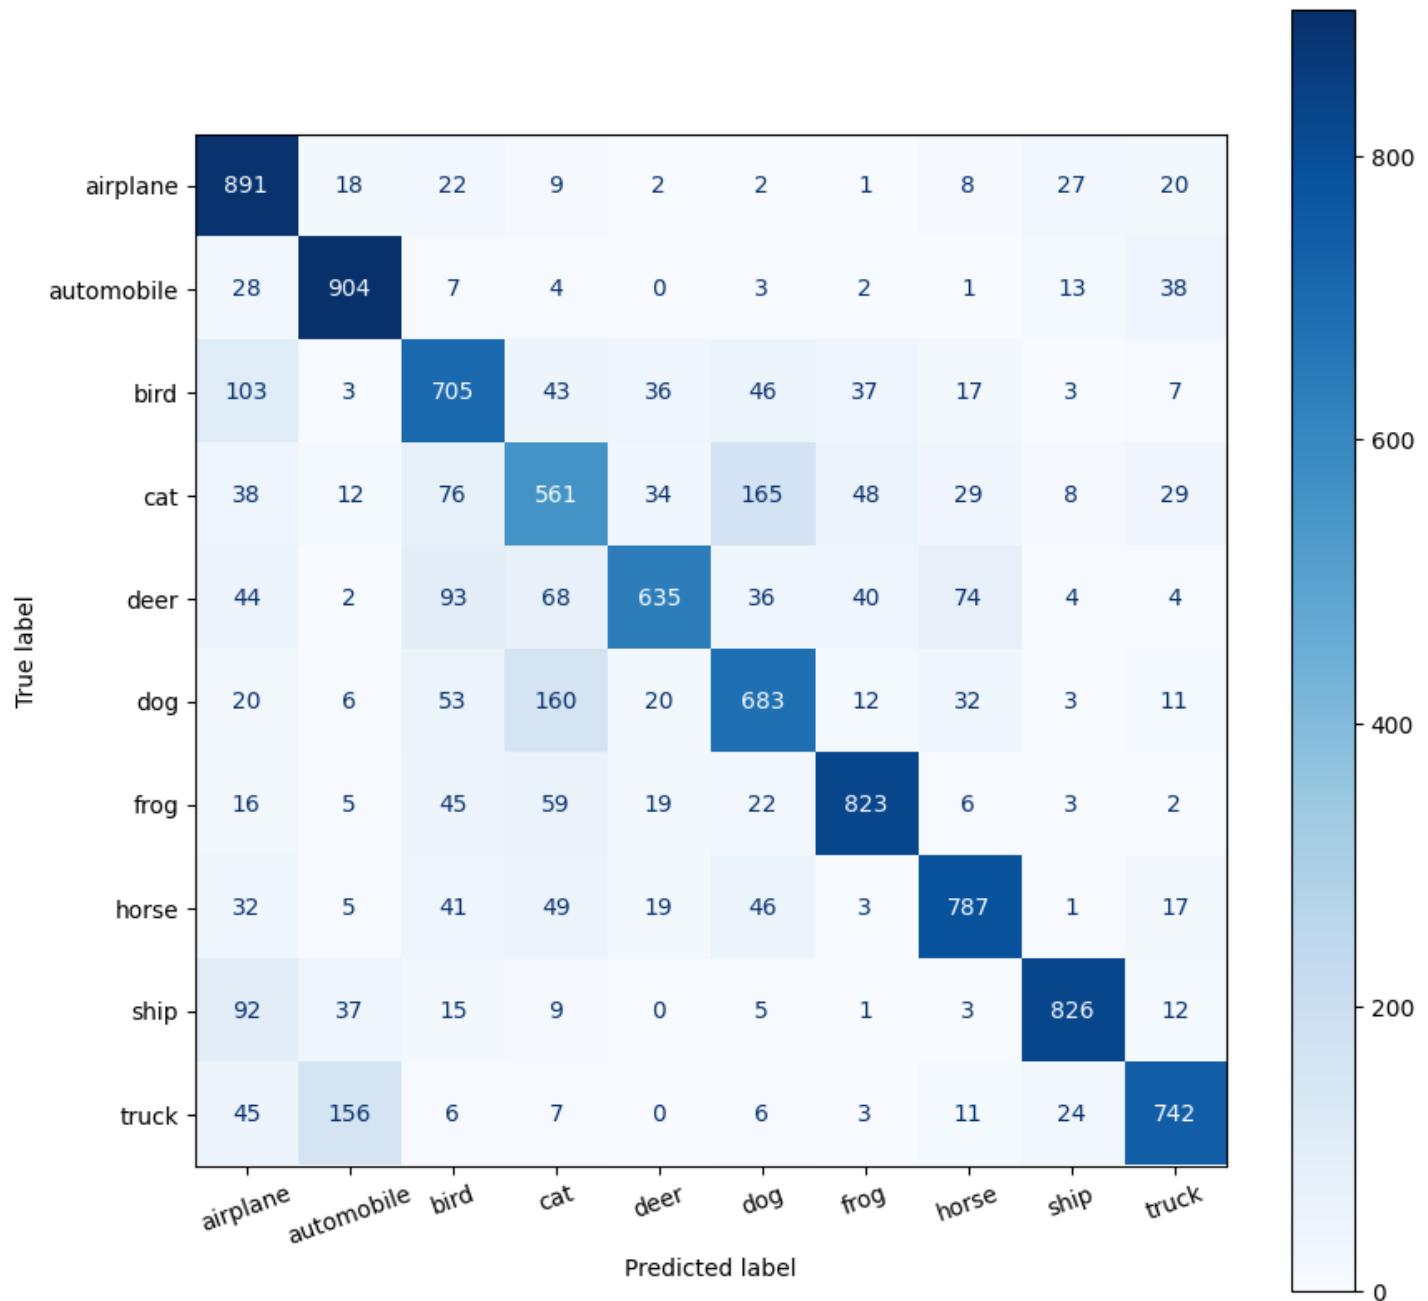
```
labels, probs = get_predictions(model, test_iterator, device)
```

```
pred_labels = torch.argmax(probs, 1)
```

```
def plot_confusion_matrix(labels, pred_labels, classes):

    fig = plt.figure(figsize = (10, 10));
    ax = fig.add_subplot(1, 1, 1);
    cm = confusion_matrix(labels.cpu(), pred_labels.cpu());
    cm = ConfusionMatrixDisplay(cm, display_labels = classes);
    cm.plot(values_format = 'd', cmap = 'Blues', ax = ax)
    plt.xticks(rotation = 20)
```

```
plot_confusion_matrix(labels, pred_labels, classes)
```



Conclusion

That's it! As a side project (this is not for credit and won't be graded), feel free to play around with different design choices that you made while building this network.

- Whether or not to normalize the color channels in the input.
- The learning rate parameter in Adam.
- The batch size.
- The number of training epochs.
- (and if you are feeling brave -- the AlexNet architecture itself.)

TorchVision Object Detection Finetuning Tutorial

Defining the Dataset

```
import os
import torch

from torchvision.io import read_image
from torchvision.ops.boxes import masks_to_boxes
from torchvision import tv_tensors
from torchvision.transforms.v2 import functional as F

import warnings
warnings.filterwarnings("ignore")

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root, "PNGImages")))))
        self.masks = list(sorted(os.listdir(os.path.join(root, "PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
        img_path = os.path.join(self.root, "PNGImages", self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks", self.masks[idx])
        img = read_image(img_path)
        mask = read_image(mask_path)
        # instances are encoded as different colors
        obj_ids = torch.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]
        num_objs = len(obj_ids)

        # split the color-encoded mask into a set
        # of binary masks
        masks = (mask == obj_ids[:, None, None]).to(dtype=torch.uint8)

        # get bounding box coordinates for each mask
        boxes = masks_to_boxes(masks)

        # there is only one class
        labels = torch.ones((num_objs,), dtype=torch.int64)

        image_id = idx
```

```

area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0])
# suppose all instances are not crowd
iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

# Wrap sample and targets into torchvision tv_tensors:
img = tv_tensors.Image(img)

target = {}
target["boxes"] = tv_tensors.BoundingBoxes(boxes, format="XYXY",
canvas_size=F.get_size(img))
target["masks"] = tv_tensors.Mask(masks)
target["labels"] = labels
target["image_id"] = image_id
target["area"] = area
target["iscrowd"] = iscrowd

if self.transforms is not None:
    img, target = self.transforms(img, target)

return img, target

def __len__(self):
    return len(self.imgs)

```

1 - Finetuning from a pretrained model

```

import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor


def get_model_instance_segmentation(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model = torchvision.models.detection.maskrcnn_resnet50_fpn(weights="DEFAULT")

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask = model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256
    # and replace the mask predictor with a new one
    model.roi_heads.mask_predictor = MaskRCNNPredictor(
        in_features_mask,
        hidden_layer,
        num_classes

```

```

        )

    return model

from torchvision.transforms import v2 as T

def get_transform(train):
    transforms = []
    if train:
        transforms.append(T.RandomHorizontalFlip(0.5))
    transforms.append(T.ToDtype(torch.float, scale=True))
    transforms.append(T.ToPureTensor())
    return T.Compose(transforms)

import utils
from engine import train_one_epoch, evaluate

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2
# use our dataset and defined transformations
dataset = PennFudanDataset('PennFudanPed', get_transform(train=True))
dataset_test = PennFudanDataset('PennFudanPed', get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=0,
    collate_fn=utils.collate_fn
)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test,
    batch_size=1,
    shuffle=False,
    num_workers=0,
)

```

```

        collate_fn=utils.collate_fn
    )

# get the model using our helper function
model = get_model_instance_segmentation(num_classes)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs, modify to 10
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

print("That's it!")

```

```

Epoch: [0]  [ 0/60]  eta: 0:07:02  lr: 0.000090  loss: 5.3390 (5.3390)  loss_classifier:
1.0105 (1.0105)  loss_box_reg: 0.1436 (0.1436)  loss_mask: 4.1724 (4.1724)  loss_objectness:
0.0037 (0.0037)  loss_rpn_box_reg: 0.0088 (0.0088)  time: 7.0445  data: 0.0234  max mem: 2258
Epoch: [0]  [10/60]  eta: 0:00:57  lr: 0.000936  loss: 1.8463 (2.7963)  loss_classifier:
0.5531 (0.5893)  loss_box_reg: 0.2991 (0.2884)  loss_mask: 1.2507 (1.8919)  loss_objectness:
0.0214 (0.0224)  loss_rpn_box_reg: 0.0041 (0.0043)  time: 1.1425  data: 0.0330  max mem: 2768
Epoch: [0]  [20/60]  eta: 0:00:34  lr: 0.001783  loss: 1.1690 (1.7923)  loss_classifier:
0.2832 (0.3868)  loss_box_reg: 0.2921 (0.2685)  loss_mask: 0.4370 (1.1116)  loss_objectness:
0.0203 (0.0202)  loss_rpn_box_reg: 0.0041 (0.0051)  time: 0.5459  data: 0.0343  max mem: 2768
Epoch: [0]  [30/60]  eta: 0:00:22  lr: 0.002629  loss: 0.5333 (1.3968)  loss_classifier:

```

```

0.0836 (0.2878) loss_box_reg: 0.1943 (0.2556) loss_mask: 0.2266 (0.8300) loss_objectness:
0.0108 (0.0169) loss_rpn_box_reg: 0.0054 (0.0065) time: 0.5423 data: 0.0350 max mem: 2768
Epoch: [0] [40/60] eta: 0:00:14 lr: 0.003476 loss: 0.5064 (1.1767) loss_classifier:
0.0605 (0.2309) loss_box_reg: 0.1882 (0.2449) loss_mask: 0.2233 (0.6801) loss_objectness:
0.0070 (0.0142) loss_rpn_box_reg: 0.0060 (0.0066) time: 0.5498 data: 0.0359 max mem: 2768
Epoch: [0] [50/60] eta: 0:00:06 lr: 0.004323 loss: 0.4367 (1.0277) loss_classifier:
0.0496 (0.1949) loss_box_reg: 0.1796 (0.2326) loss_mask: 0.1884 (0.5817) loss_objectness:
0.0040 (0.0121) loss_rpn_box_reg: 0.0051 (0.0064) time: 0.5473 data: 0.0349 max mem: 2857
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3930 (0.9301) loss_classifier:
0.0411 (0.1716) loss_box_reg: 0.1604 (0.2184) loss_mask: 0.1720 (0.5232) loss_objectness:
0.0018 (0.0105) loss_rpn_box_reg: 0.0050 (0.0063) time: 0.5446 data: 0.0337 max mem: 2857
Epoch: [0] Total time: 0:00:39 (0.6535 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1604 (0.1604) evaluator_time: 0.0114 (0.0114)
time: 0.1811 data: 0.0086 max mem: 2857
Test: [49/50] eta: 0:00:00 model_time: 0.0941 (0.1085) evaluator_time: 0.0055 (0.0091)
time: 0.1253 data: 0.0124 max mem: 2857
Test: Total time: 0:00:06 (0.1311 s / it)
Averaged stats: model_time: 0.0941 (0.1085) evaluator_time: 0.0055 (0.0091)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.619
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.990
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.749
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.425
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.524
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.627
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.263
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.675
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.675
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.650
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.644
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.677
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.667
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.984
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.870
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.300
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.547
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.676
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.276
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.704
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.709
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.600
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.744

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.708
Epoch: [1] [ 0/60] eta: 0:00:27 lr: 0.005000 loss: 0.2231 (0.2231) loss_classifier:
0.0209 (0.0209) loss_box_reg: 0.0553 (0.0553) loss_mask: 0.1395 (0.1395) loss_objectness:
0.0016 (0.0016) loss_rpn_box_reg: 0.0058 (0.0058) time: 0.4534 data: 0.0213 max mem: 2857
Epoch: [1] [10/60] eta: 0:00:26 lr: 0.005000 loss: 0.2564 (0.2813) loss_classifier:
0.0267 (0.0389) loss_box_reg: 0.0777 (0.0952) loss_mask: 0.1418 (0.1422) loss_objectness:
0.0010 (0.0014) loss_rpn_box_reg: 0.0026 (0.0036) time: 0.5397 data: 0.0259 max mem: 2884
Epoch: [1] [20/60] eta: 0:00:22 lr: 0.005000 loss: 0.2787 (0.2950) loss_classifier:
0.0319 (0.0378) loss_box_reg: 0.0879 (0.1015) loss_mask: 0.1418 (0.1493) loss_objectness:
0.0012 (0.0018) loss_rpn_box_reg: 0.0030 (0.0047) time: 0.5784 data: 0.0385 max mem: 2884
Epoch: [1] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.3099 (0.3159) loss_classifier:
0.0355 (0.0404) loss_box_reg: 0.1161 (0.1106) loss_mask: 0.1509 (0.1572) loss_objectness:
0.0016 (0.0022) loss_rpn_box_reg: 0.0061 (0.0055) time: 0.6499 data: 0.0586 max mem: 2884
Epoch: [1] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.3249 (0.3128) loss_classifier:
0.0407 (0.0409) loss_box_reg: 0.1024 (0.1063) loss_mask: 0.1509 (0.1576) loss_objectness:
0.0011 (0.0021) loss_rpn_box_reg: 0.0061 (0.0058) time: 0.6338 data: 0.0513 max mem: 2884
Epoch: [1] [50/60] eta: 0:00:06 lr: 0.005000 loss: 0.2723 (0.3133) loss_classifier:
0.0473 (0.0432) loss_box_reg: 0.0887 (0.1065) loss_mask: 0.1439 (0.1558) loss_objectness:
0.0010 (0.0021) loss_rpn_box_reg: 0.0033 (0.0056) time: 0.6041 data: 0.0381 max mem: 2948
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2641 (0.3072) loss_classifier:
0.0395 (0.0423) loss_box_reg: 0.0795 (0.1022) loss_mask: 0.1405 (0.1555) loss_objectness:
0.0008 (0.0019) loss_rpn_box_reg: 0.0031 (0.0052) time: 0.5998 data: 0.0361 max mem: 2948
Epoch: [1] Total time: 0:00:36 (0.6005 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:08 model_time: 0.1473 (0.1473) evaluator_time: 0.0223 (0.0223)
time: 0.1800 data: 0.0097 max mem: 2948
Test: [49/50] eta: 0:00:00 model_time: 0.1074 (0.1131) evaluator_time: 0.0051 (0.0091)
time: 0.1395 data: 0.0187 max mem: 2948
Test: Total time: 0:00:07 (0.1422 s / it)
Averaged stats: model_time: 0.1074 (0.1131) evaluator_time: 0.0051 (0.0091)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.791
Average Precision (AP) @[ IoU=0.50           | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75           | area= all | maxDets=100 ] = 0.923
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.533
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.590
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.801
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.334
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.834
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.834
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.800
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.839
IoU metric: segm

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.733
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.947
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.300
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.586
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.742
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.310
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.773
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.773
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.600
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.775
Epoch: [2] [ 0/60] eta: 0:00:33 lr: 0.005000 loss: 0.2981 (0.2981) loss_classifier:
0.0507 (0.0507) loss_box_reg: 0.0932 (0.0932) loss_mask: 0.1490 (0.1490) loss_objectness:
0.0018 (0.0018) loss_rpn_box_reg: 0.0036 (0.0036) time: 0.5648 data: 0.0228 max mem: 2948
Epoch: [2] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2350 (0.2414) loss_classifier:
0.0404 (0.0398) loss_box_reg: 0.0616 (0.0669) loss_mask: 0.1205 (0.1286) loss_objectness:
0.0008 (0.0022) loss_rpn_box_reg: 0.0022 (0.0039) time: 0.5809 data: 0.0379 max mem: 2948
Epoch: [2] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2350 (0.2525) loss_classifier:
0.0371 (0.0375) loss_box_reg: 0.0577 (0.0703) loss_mask: 0.1294 (0.1382) loss_objectness:
0.0008 (0.0018) loss_rpn_box_reg: 0.0028 (0.0047) time: 0.5999 data: 0.0384 max mem: 2949
Epoch: [2] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2196 (0.2391) loss_classifier:
0.0282 (0.0336) loss_box_reg: 0.0498 (0.0626) loss_mask: 0.1400 (0.1365) loss_objectness:
0.0005 (0.0019) loss_rpn_box_reg: 0.0022 (0.0045) time: 0.5980 data: 0.0313 max mem: 2949
Epoch: [2] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.2196 (0.2449) loss_classifier:
0.0285 (0.0343) loss_box_reg: 0.0507 (0.0669) loss_mask: 0.1400 (0.1373) loss_objectness:
0.0005 (0.0017) loss_rpn_box_reg: 0.0022 (0.0047) time: 0.6028 data: 0.0285 max mem: 3130
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2274 (0.2408) loss_classifier:
0.0288 (0.0323) loss_box_reg: 0.0716 (0.0667) loss_mask: 0.1338 (0.1354) loss_objectness:
0.0008 (0.0018) loss_rpn_box_reg: 0.0041 (0.0045) time: 0.5979 data: 0.0266 max mem: 3130
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2327 (0.2408) loss_classifier:
0.0240 (0.0322) loss_box_reg: 0.0654 (0.0676) loss_mask: 0.1318 (0.1350) loss_objectness:
0.0006 (0.0016) loss_rpn_box_reg: 0.0034 (0.0044) time: 0.5847 data: 0.0278 max mem: 3130
Epoch: [2] Total time: 0:00:35 (0.5940 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:07 model_time: 0.1406 (0.1406) evaluator_time: 0.0030 (0.0030)
time: 0.1527 data: 0.0084 max mem: 3130
Test: [49/50] eta: 0:00:00 model_time: 0.1040 (0.1116) evaluator_time: 0.0055 (0.0088)
time: 0.1499 data: 0.0195 max mem: 3130
Test: Total time: 0:00:06 (0.1377 s / it)
Averaged stats: model_time: 0.1040 (0.1116) evaluator_time: 0.0055 (0.0088)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.723
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991

```

```

Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.901
Average Precision (AP) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.326
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.566
Average Precision (AP) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.733
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.308
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.768
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.768
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.650
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.722
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.773
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.735
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.982
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.912
Average Precision (AP) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.240
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.530
Average Precision (AP) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.749
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.314
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.773
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.775
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.600
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.779
Epoch: [3] [ 0/60] eta: 0:00:38 lr: 0.000500 loss: 0.2697 (0.2697) loss_classifier:
0.0295 (0.0295) loss_box_reg: 0.1028 (0.1028) loss_mask: 0.1346 (0.1346) loss_objectness:
0.0003 (0.0003) loss_rpn_box_reg: 0.0024 (0.0024) time: 0.6364 data: 0.0373 max mem: 3130
Epoch: [3] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.2443 (0.2586) loss_classifier:
0.0288 (0.0316) loss_box_reg: 0.0740 (0.0815) loss_mask: 0.1346 (0.1399) loss_objectness:
0.0006 (0.0012) loss_rpn_box_reg: 0.0033 (0.0045) time: 0.6381 data: 0.0391 max mem: 3130
Epoch: [3] [20/60] eta: 0:00:24 lr: 0.000500 loss: 0.2143 (0.2240) loss_classifier:
0.0245 (0.0284) loss_box_reg: 0.0530 (0.0645) loss_mask: 0.1214 (0.1268) loss_objectness:
0.0005 (0.0008) loss_rpn_box_reg: 0.0026 (0.0036) time: 0.6030 data: 0.0338 max mem: 3130
Epoch: [3] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1807 (0.2138) loss_classifier:
0.0225 (0.0274) loss_box_reg: 0.0378 (0.0576) loss_mask: 0.1160 (0.1248) loss_objectness:
0.0003 (0.0007) loss_rpn_box_reg: 0.0022 (0.0033) time: 0.5729 data: 0.0283 max mem: 3130
Epoch: [3] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1791 (0.2100) loss_classifier:
0.0227 (0.0276) loss_box_reg: 0.0350 (0.0544) loss_mask: 0.1156 (0.1243) loss_objectness:
0.0004 (0.0007) loss_rpn_box_reg: 0.0016 (0.0031) time: 0.6032 data: 0.0302 max mem: 3134
Epoch: [3] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1774 (0.2055) loss_classifier:
0.0219 (0.0271) loss_box_reg: 0.0295 (0.0516) loss_mask: 0.1086 (0.1230) loss_objectness:
0.0003 (0.0008) loss_rpn_box_reg: 0.0012 (0.0029) time: 0.5961 data: 0.0309 max mem: 3134
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1740 (0.2028) loss_classifier:
0.0207 (0.0270) loss_box_reg: 0.0321 (0.0505) loss_mask: 0.1086 (0.1215) loss_objectness:
0.0004 (0.0009) loss_rpn_box_reg: 0.0021 (0.0029) time: 0.5913 data: 0.0326 max mem: 3134
Epoch: [3] Total time: 0:00:35 (0.5932 s / it)
creating index...
index created!

```

```
Test: [ 0/50] eta: 0:00:07 model_time: 0.1401 (0.1401) evaluator_time: 0.0027 (0.0027)
```

```
time: 0.1520  data: 0.0085  max mem: 3134
Test: [49/50]  eta: 0:00:00  model_time: 0.0995 (0.1062)  evaluator_time: 0.0033 (0.0056)
time: 0.1283  data: 0.0141  max mem: 3134
Test: Total time: 0:00:06 (0.1259 s / it)
Averaged stats: model_time: 0.0995 (0.1062)  evaluator_time: 0.0033 (0.0056)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.02s).

IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.818
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.945
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.552
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.675
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.828
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.345
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.855
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.855
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.700
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.844
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.858

IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.745
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.934
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.296
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.564
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.315
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.650
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.778
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.790

Epoch: [4]  [ 0/60]  eta: 0:00:45  lr: 0.000500  loss: 0.1845 (0.1845)  loss_classifier: 0.0239 (0.0239)  loss_box_reg: 0.0420 (0.0420)  loss_mask: 0.1169 (0.1169)  loss_objectness: 0.0003 (0.0003)  loss_rpn_box_reg: 0.0014 (0.0014)  time: 0.7579  data: 0.0463  max mem: 3134
Epoch: [4]  [10/60]  eta: 0:00:34  lr: 0.000500  loss: 0.1845 (0.1758)  loss_classifier: 0.0239 (0.0209)  loss_box_reg: 0.0420 (0.0374)  loss_mask: 0.1091 (0.1143)  loss_objectness: 0.0003 (0.0009)  loss_rpn_box_reg: 0.0027 (0.0023)  time: 0.6819  data: 0.0567  max mem: 3134
Epoch: [4]  [20/60]  eta: 0:00:25  lr: 0.000500  loss: 0.1957 (0.1950)  loss_classifier: 0.0262 (0.0271)  loss_box_reg: 0.0411 (0.0446)  loss_mask: 0.1121 (0.1196)  loss_objectness: 0.0004 (0.0012)  loss_rpn_box_reg: 0.0022 (0.0024)  time: 0.6363  data: 0.0437  max mem: 3134
Epoch: [4]  [30/60]  eta: 0:00:19  lr: 0.000500  loss: 0.1831 (0.1911)  loss_classifier: 0.0243 (0.0265)  loss_box_reg: 0.0411 (0.0443)  loss_mask: 0.1086 (0.1168)  loss_objectness: 0.0006 (0.0011)  loss_rpn_box_reg: 0.0017 (0.0024)  time: 0.6393  data: 0.0398  max mem: 3210

Epoch: [4]  [40/60]  eta: 0:00:12  lr: 0.000500  loss: 0.1830 (0.1968)  loss_classifier: 0.0220 (0.0275)  loss_box_reg: 0.0424 (0.0450)  loss_mask: 0.1000 (0.1004)  loss_objectness: 0.0005 (0.0010)  loss_rpn_box_reg: 0.0016 (0.0023)  time: 0.6363  data: 0.0398  max mem: 3210
```

```

0.0230 (0.0215) loss_box_reg: 0.0434 (0.0452) loss_mask: 0.1099 (0.1204) loss_objectness:
0.0006 (0.0009) loss_rpn_box_reg: 0.0017 (0.0027) time: 0.6220 data: 0.0386 max mem: 3210
Epoch: [4] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.2034 (0.1979) loss_classifier:
0.0280 (0.0281) loss_box_reg: 0.0434 (0.0455) loss_mask: 0.1244 (0.1206) loss_objectness:
0.0005 (0.0009) loss_rpn_box_reg: 0.0036 (0.0028) time: 0.5994 data: 0.0386 max mem: 3210
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1706 (0.1935) loss_classifier:
0.0244 (0.0269) loss_box_reg: 0.0338 (0.0437) loss_mask: 0.1080 (0.1192) loss_objectness:
0.0005 (0.0009) loss_rpn_box_reg: 0.0031 (0.0028) time: 0.6260 data: 0.0436 max mem: 3210
Epoch: [4] Total time: 0:00:37 (0.6332 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:07 model_time: 0.1430 (0.1430) evaluator_time: 0.0039 (0.0039)
time: 0.1595 data: 0.0119 max mem: 3210
Test: [49/50] eta: 0:00:00 model_time: 0.1010 (0.1149) evaluator_time: 0.0036 (0.0076)
time: 0.1397 data: 0.0166 max mem: 3210
Test: Total time: 0:00:06 (0.1398 s / it)
Averaged stats: model_time: 0.1010 (0.1149) evaluator_time: 0.0036 (0.0076)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).

IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.823
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.943
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.350
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.672
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.834
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.345
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.861
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.861
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.700
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.833
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.866

IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.751
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.932
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.325
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.563
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.762
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.315
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.650
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.800
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.794

Epoch: [5] [ 0/60] eta: 0:00:29 lr: 0.000500 loss: 0.2033 (0.2033) loss_classifier:
0.0236 (0.0236) loss_box_reg: 0.0261 (0.0261) loss_mask: 0.1412 (0.1412) loss_objectness:

```

```

loss_box_reg: 0.0001 (0.0001) loss_mask: 0.1415 (0.1415) loss_objectness:
0.0002 (0.0002) loss_rpn_box_reg: 0.0021 (0.0021) time: 0.4895 data: 0.0308 max mem: 3210
Epoch: [5] [10/60] eta: 0:00:28 lr: 0.000500 loss: 0.1743 (0.1834) loss_classifier:
0.0233 (0.0246) loss_box_reg: 0.0330 (0.0387) loss_mask: 0.1176 (0.1175) loss_objectness:
0.0003 (0.0003) loss_rpn_box_reg: 0.0011 (0.0023) time: 0.5638 data: 0.0297 max mem: 3210
Epoch: [5] [20/60] eta: 0:00:22 lr: 0.000500 loss: 0.1743 (0.1912) loss_classifier:
0.0242 (0.0278) loss_box_reg: 0.0345 (0.0416) loss_mask: 0.1176 (0.1183) loss_objectness:
0.0004 (0.0007) loss_rpn_box_reg: 0.0018 (0.0028) time: 0.5777 data: 0.0296 max mem: 3210
Epoch: [5] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1816 (0.1929) loss_classifier:
0.0242 (0.0270) loss_box_reg: 0.0395 (0.0437) loss_mask: 0.1194 (0.1182) loss_objectness:
0.0005 (0.0010) loss_rpn_box_reg: 0.0024 (0.0030) time: 0.5798 data: 0.0298 max mem: 3210
Epoch: [5] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1713 (0.1865) loss_classifier:
0.0234 (0.0267) loss_box_reg: 0.0378 (0.0413) loss_mask: 0.1095 (0.1150) loss_objectness:
0.0004 (0.0008) loss_rpn_box_reg: 0.0017 (0.0027) time: 0.5919 data: 0.0281 max mem: 3500
Epoch: [5] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1708 (0.1864) loss_classifier:
0.0233 (0.0261) loss_box_reg: 0.0330 (0.0409) loss_mask: 0.1058 (0.1158) loss_objectness:
0.0003 (0.0008) loss_rpn_box_reg: 0.0017 (0.0027) time: 0.6047 data: 0.0297 max mem: 3500
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1711 (0.1862) loss_classifier:
0.0198 (0.0256) loss_box_reg: 0.0343 (0.0405) loss_mask: 0.1122 (0.1165) loss_objectness:
0.0003 (0.0008) loss_rpn_box_reg: 0.0026 (0.0028) time: 0.5860 data: 0.0284 max mem: 3500
Epoch: [5] Total time: 0:00:35 (0.5854 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:07 model_time: 0.1435 (0.1435) evaluator_time: 0.0025 (0.0025)
time: 0.1561 data: 0.0093 max mem: 3500
Test: [49/50] eta: 0:00:00 model_time: 0.0991 (0.1052) evaluator_time: 0.0031 (0.0057)
time: 0.1220 data: 0.0134 max mem: 3500
Test: Total time: 0:00:06 (0.1260 s / it)
Averaged stats: model_time: 0.0991 (0.1052) evaluator_time: 0.0031 (0.0057)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.830
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.635
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.652
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.840
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.355
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.869
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.869
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.822
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.875
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.752
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.922

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.325
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.578
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.761
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.316
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.791
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.793
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.650
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.822
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.793
Epoch: [6] [ 0/60] eta: 0:00:41 lr: 0.000050 loss: 0.1881 (0.1881) loss_classifier:
0.0245 (0.0245) loss_box_reg: 0.0411 (0.0411) loss_mask: 0.1181 (0.1181) loss_objectness:
0.0002 (0.0002) loss_rpn_box_reg: 0.0043 (0.0043) time: 0.6954 data: 0.0222 max mem: 3500
Epoch: [6] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1881 (0.1951) loss_classifier:
0.0245 (0.0262) loss_box_reg: 0.0335 (0.0410) loss_mask: 0.1181 (0.1232) loss_objectness:
0.0002 (0.0018) loss_rpn_box_reg: 0.0011 (0.0030) time: 0.5858 data: 0.0288 max mem: 3500
Epoch: [6] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1755 (0.2008) loss_classifier:
0.0256 (0.0273) loss_box_reg: 0.0322 (0.0451) loss_mask: 0.1131 (0.1236) loss_objectness:
0.0002 (0.0014) loss_rpn_box_reg: 0.0032 (0.0034) time: 0.5979 data: 0.0331 max mem: 3500
Epoch: [6] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1677 (0.1868) loss_classifier:
0.0215 (0.0247) loss_box_reg: 0.0289 (0.0393) loss_mask: 0.1119 (0.1188) loss_objectness:
0.0003 (0.0012) loss_rpn_box_reg: 0.0019 (0.0028) time: 0.5783 data: 0.0300 max mem: 3500
Epoch: [6] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1590 (0.1862) loss_classifier:
0.0209 (0.0248) loss_box_reg: 0.0299 (0.0390) loss_mask: 0.1084 (0.1186) loss_objectness:
0.0003 (0.0009) loss_rpn_box_reg: 0.0016 (0.0028) time: 0.5592 data: 0.0272 max mem: 3500
Epoch: [6] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1878 (0.1903) loss_classifier:
0.0219 (0.0255) loss_box_reg: 0.0304 (0.0406) loss_mask: 0.1092 (0.1203) loss_objectness:
0.0004 (0.0011) loss_rpn_box_reg: 0.0025 (0.0029) time: 0.5829 data: 0.0290 max mem: 3500
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1732 (0.1877) loss_classifier:
0.0232 (0.0254) loss_box_reg: 0.0351 (0.0401) loss_mask: 0.1071 (0.1185) loss_objectness:
0.0003 (0.0009) loss_rpn_box_reg: 0.0019 (0.0027) time: 0.6117 data: 0.0417 max mem: 3500
Epoch: [6] Total time: 0:00:35 (0.5897 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:07 model_time: 0.1407 (0.1407) evaluator_time: 0.0024 (0.0024)
time: 0.1521 data: 0.0083 max mem: 3500
Test: [49/50] eta: 0:00:00 model_time: 0.0979 (0.1042) evaluator_time: 0.0032 (0.0049)
time: 0.1201 data: 0.0126 max mem: 3500
Test: Total time: 0:00:06 (0.1225 s / it)
Averaged stats: model_time: 0.0979 (0.1042) evaluator_time: 0.0032 (0.0049)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.831
Average Precision (AP) @[ IoU=0.50          | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75          | area= all | maxDets=100 ] = 0.943
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.635
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.667

```

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.842
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.352
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.872
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.872
 Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.750
 Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.833
 Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.876

IoU metric: segm

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.753
 Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.991
 Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.940
 Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.325
 Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.567
 Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.763
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.315
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.791
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.793
 Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.650
 Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.800
 Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.795

Epoch: [7] [0/60] eta: 0:00:41 lr: 0.000050 loss: 0.2865 (0.2865) loss_classifier: 0.0411 (0.0411) loss_box_reg: 0.0862 (0.0862) loss_mask: 0.1506 (0.1506) loss_objectness: 0.0004 (0.0004) loss_rpn_box_reg: 0.0083 (0.0083) time: 0.6939 data: 0.0432 max mem: 3500

Epoch: [7] [10/60] eta: 0:00:30 lr: 0.000050 loss: 0.1941 (0.2085) loss_classifier: 0.0256 (0.0280) loss_box_reg: 0.0464 (0.0523) loss_mask: 0.1192 (0.1238) loss_objectness: 0.0002 (0.0007) loss_rpn_box_reg: 0.0021 (0.0037) time: 0.6102 data: 0.0327 max mem: 3500

Epoch: [7] [20/60] eta: 0:00:23 lr: 0.000050 loss: 0.1610 (0.1853) loss_classifier: 0.0190 (0.0242) loss_box_reg: 0.0286 (0.0425) loss_mask: 0.1046 (0.1150) loss_objectness: 0.0003 (0.0007) loss_rpn_box_reg: 0.0015 (0.0030) time: 0.5880 data: 0.0289 max mem: 3500

Epoch: [7] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1633 (0.1897) loss_classifier: 0.0219 (0.0251) loss_box_reg: 0.0300 (0.0434) loss_mask: 0.1019 (0.1178) loss_objectness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0015 (0.0028) time: 0.5838 data: 0.0270 max mem: 3500

Epoch: [7] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1785 (0.1851) loss_classifier: 0.0243 (0.0249) loss_box_reg: 0.0357 (0.0406) loss_mask: 0.1096 (0.1165) loss_objectness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0019 (0.0026) time: 0.5719 data: 0.0275 max mem: 3500

Epoch: [7] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1785 (0.1863) loss_classifier: 0.0243 (0.0248) loss_box_reg: 0.0377 (0.0407) loss_mask: 0.1096 (0.1175) loss_objectness: 0.0003 (0.0005) loss_rpn_box_reg: 0.0020 (0.0027) time: 0.5665 data: 0.0281 max mem: 3500

Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1808 (0.1837) loss_classifier: 0.0232 (0.0239) loss_box_reg: 0.0356 (0.0394) loss_mask: 0.1155 (0.1173) loss_objectness: 0.0002 (0.0005) loss_rpn_box_reg: 0.0024 (0.0027) time: 0.5784 data: 0.0276 max mem: 3500

Epoch: [7] Total time: 0:00:34 (0.5818 s / it)

creating index...

index created!

Test: [0/50] eta: 0:00:08 model_time: 0.1437 (0.1437) evaluator_time: 0.0040 (0.0040) time: 0.1603 data: 0.0119 max mem: 3500

Test: [49/50] eta: 0:00:00 model_time: 0.0981 (0.1066) evaluator_time: 0.0032 (0.0056) time: 0.1211 data: 0.0127 max mem: 3500

```
Test: Total time: 0:00:06 (0.1269 s / it)
Averaged stats: model_time: 0.0981 (0.1066) evaluator_time: 0.0032 (0.0056)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.826
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.943
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.635
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.667
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.836
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.350
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.868
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.868
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.833
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.872
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.750
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.922
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.325
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.555
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.762
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.315
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.650
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.794
Epoch: [8] [ 0/60] eta: 0:00:38 lr: 0.000050 loss: 0.1496 (0.1496) loss_classifier: 0.0159 (0.0159) loss_box_reg: 0.0264 (0.0264) loss_mask: 0.1060 (0.1060) loss_objectness: 0.0001 (0.0001) loss_rpn_box_reg: 0.0013 (0.0013) time: 0.6364 data: 0.0282 max mem: 3500
Epoch: [8] [10/60] eta: 0:00:31 lr: 0.000050 loss: 0.1572 (0.1686) loss_classifier: 0.0257 (0.0236) loss_box_reg: 0.0264 (0.0314) loss_mask: 0.1060 (0.1116) loss_objectness: 0.0002 (0.0003) loss_rpn_box_reg: 0.0013 (0.0018) time: 0.6390 data: 0.0344 max mem: 3500
Epoch: [8] [20/60] eta: 0:00:25 lr: 0.000050 loss: 0.1640 (0.1732) loss_classifier: 0.0263 (0.0256) loss_box_reg: 0.0368 (0.0341) loss_mask: 0.1095 (0.1111) loss_objectness: 0.0002 (0.0007) loss_rpn_box_reg: 0.0013 (0.0017) time: 0.6406 data: 0.0463 max mem: 3500
Epoch: [8] [30/60] eta: 0:00:18 lr: 0.000050 loss: 0.1682 (0.1757) loss_classifier: 0.0248 (0.0261) loss_box_reg: 0.0359 (0.0342) loss_mask: 0.1121 (0.1127) loss_objectness: 0.0004 (0.0006) loss_rpn_box_reg: 0.0016 (0.0021) time: 0.6276 data: 0.0432 max mem: 3500
Epoch: [8] [40/60] eta: 0:00:12 lr: 0.000050 loss: 0.1817 (0.1818) loss_classifier: 0.0250 (0.0265) loss_box_reg: 0.0359 (0.0374) loss_mask: 0.1151 (0.1147) loss_objectness: 0.0004 (0.0007) loss_rpn_box_reg: 0.0026 (0.0025) time: 0.6003 data: 0.0287 max mem: 3500
Epoch: [8] [50/60] eta: 0:00:06 lr: 0.000050 loss: 0.2130 (0.1853) loss_classifier: 0.0259 (0.0264) loss_box_reg: 0.0491 (0.0393) loss_mask: 0.1160 (0.1162) loss_objectness:
```

```

0.0004 (0.0009) loss_rpn_box_reg: 0.0025 (0.0026) time: 0.5974 data: 0.0297 max mem: 3500
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1719 (0.1852) loss_classifier:
0.0240 (0.0257) loss_box_reg: 0.0347 (0.0389) loss_mask: 0.1119 (0.1170) loss_objectness:
0.0003 (0.0008) loss_rpn_box_reg: 0.0020 (0.0027) time: 0.5926 data: 0.0290 max mem: 3500
Epoch: [8] Total time: 0:00:36 (0.6119 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:07 model_time: 0.1387 (0.1387) evaluator_time: 0.0025 (0.0025)
time: 0.1504 data: 0.0086 max mem: 3500
Test: [49/50] eta: 0:00:00 model_time: 0.0992 (0.1070) evaluator_time: 0.0039 (0.0058)
time: 0.1278 data: 0.0155 max mem: 3500
Test: Total time: 0:00:06 (0.1281 s / it)
Averaged stats: model_time: 0.0992 (0.1070) evaluator_time: 0.0039 (0.0058)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.830
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.943
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.375
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.660
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.840
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.352
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.871
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.871
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.822
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.876
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.752
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.940
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.325
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.557
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.762
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.315
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.790
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.650
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.795
Epoch: [9] [ 0/60] eta: 0:00:28 lr: 0.000005 loss: 0.1142 (0.1142) loss_classifier:
0.0095 (0.0095) loss_box_reg: 0.0138 (0.0138) loss_mask: 0.0905 (0.0905) loss_objectness:
0.0000 (0.0000) loss_rpn_box_reg: 0.0004 (0.0004) time: 0.4719 data: 0.0173 max mem: 3500

Epoch: [9] [10/60] eta: 0:00:26 lr: 0.000005 loss: 0.1785 (0.1816) loss_classifier:
0.0195 (0.0227) loss_box_req: 0.0283 (0.0369) loss_mask: 0.1098 (0.1188) loss_objectness:

```

```

0.0003 (0.0009) loss_rpn_box_reg: 0.0013 (0.0024) time: 0.5334 data: 0.0247 max mem: 3500
Epoch: [9] [20/60] eta: 0:00:22 lr: 0.000005 loss: 0.1939 (0.1888) loss_classifier:
0.0268 (0.0267) loss_box_reg: 0.0384 (0.0405) loss_mask: 0.1098 (0.1184) loss_objectness:
0.0003 (0.0008) loss_rpn_box_reg: 0.0019 (0.0024) time: 0.5557 data: 0.0291 max mem: 3500
Epoch: [9] [30/60] eta: 0:00:17 lr: 0.000005 loss: 0.1848 (0.1884) loss_classifier:
0.0268 (0.0271) loss_box_reg: 0.0404 (0.0411) loss_mask: 0.1083 (0.1173) loss_objectness:
0.0003 (0.0007) loss_rpn_box_reg: 0.0019 (0.0023) time: 0.5903 data: 0.0302 max mem: 3500
Epoch: [9] [40/60] eta: 0:00:11 lr: 0.000005 loss: 0.1756 (0.1897) loss_classifier:
0.0242 (0.0269) loss_box_reg: 0.0373 (0.0409) loss_mask: 0.1105 (0.1187) loss_objectness:
0.0002 (0.0006) loss_rpn_box_reg: 0.0024 (0.0025) time: 0.5986 data: 0.0284 max mem: 3500
Epoch: [9] [50/60] eta: 0:00:05 lr: 0.000005 loss: 0.1688 (0.1854) loss_classifier:
0.0226 (0.0263) loss_box_reg: 0.0270 (0.0389) loss_mask: 0.1113 (0.1172) loss_objectness:
0.0002 (0.0006) loss_rpn_box_reg: 0.0014 (0.0025) time: 0.5685 data: 0.0276 max mem: 3500
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1754 (0.1848) loss_classifier:
0.0247 (0.0262) loss_box_reg: 0.0336 (0.0389) loss_mask: 0.1088 (0.1166) loss_objectness:
0.0004 (0.0006) loss_rpn_box_reg: 0.0020 (0.0026) time: 0.5861 data: 0.0289 max mem: 3500
Epoch: [9] Total time: 0:00:34 (0.5770 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:07 model_time: 0.1411 (0.1411) evaluator_time: 0.0025 (0.0025)
time: 0.1527 data: 0.0086 max mem: 3500
Test: [49/50] eta: 0:00:00 model_time: 0.0985 (0.1053) evaluator_time: 0.0033 (0.0050)
time: 0.1215 data: 0.0130 max mem: 3500
Test: Total time: 0:00:06 (0.1243 s / it)
Averaged stats: model_time: 0.0985 (0.1053) evaluator_time: 0.0033 (0.0050)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.831
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.992
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.943
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.375
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.660
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.842
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.352
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.872
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.872
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.822
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.878
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.752
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.940
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.325
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.554
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.762

```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.315
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.790
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.792
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.650
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.789
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.795
That's it!
```

```
import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes, draw_segmentation_masks

image = read_image("Beatles_-_Abbey_Road.jpg")
eval_transform = get_transform(train=False)

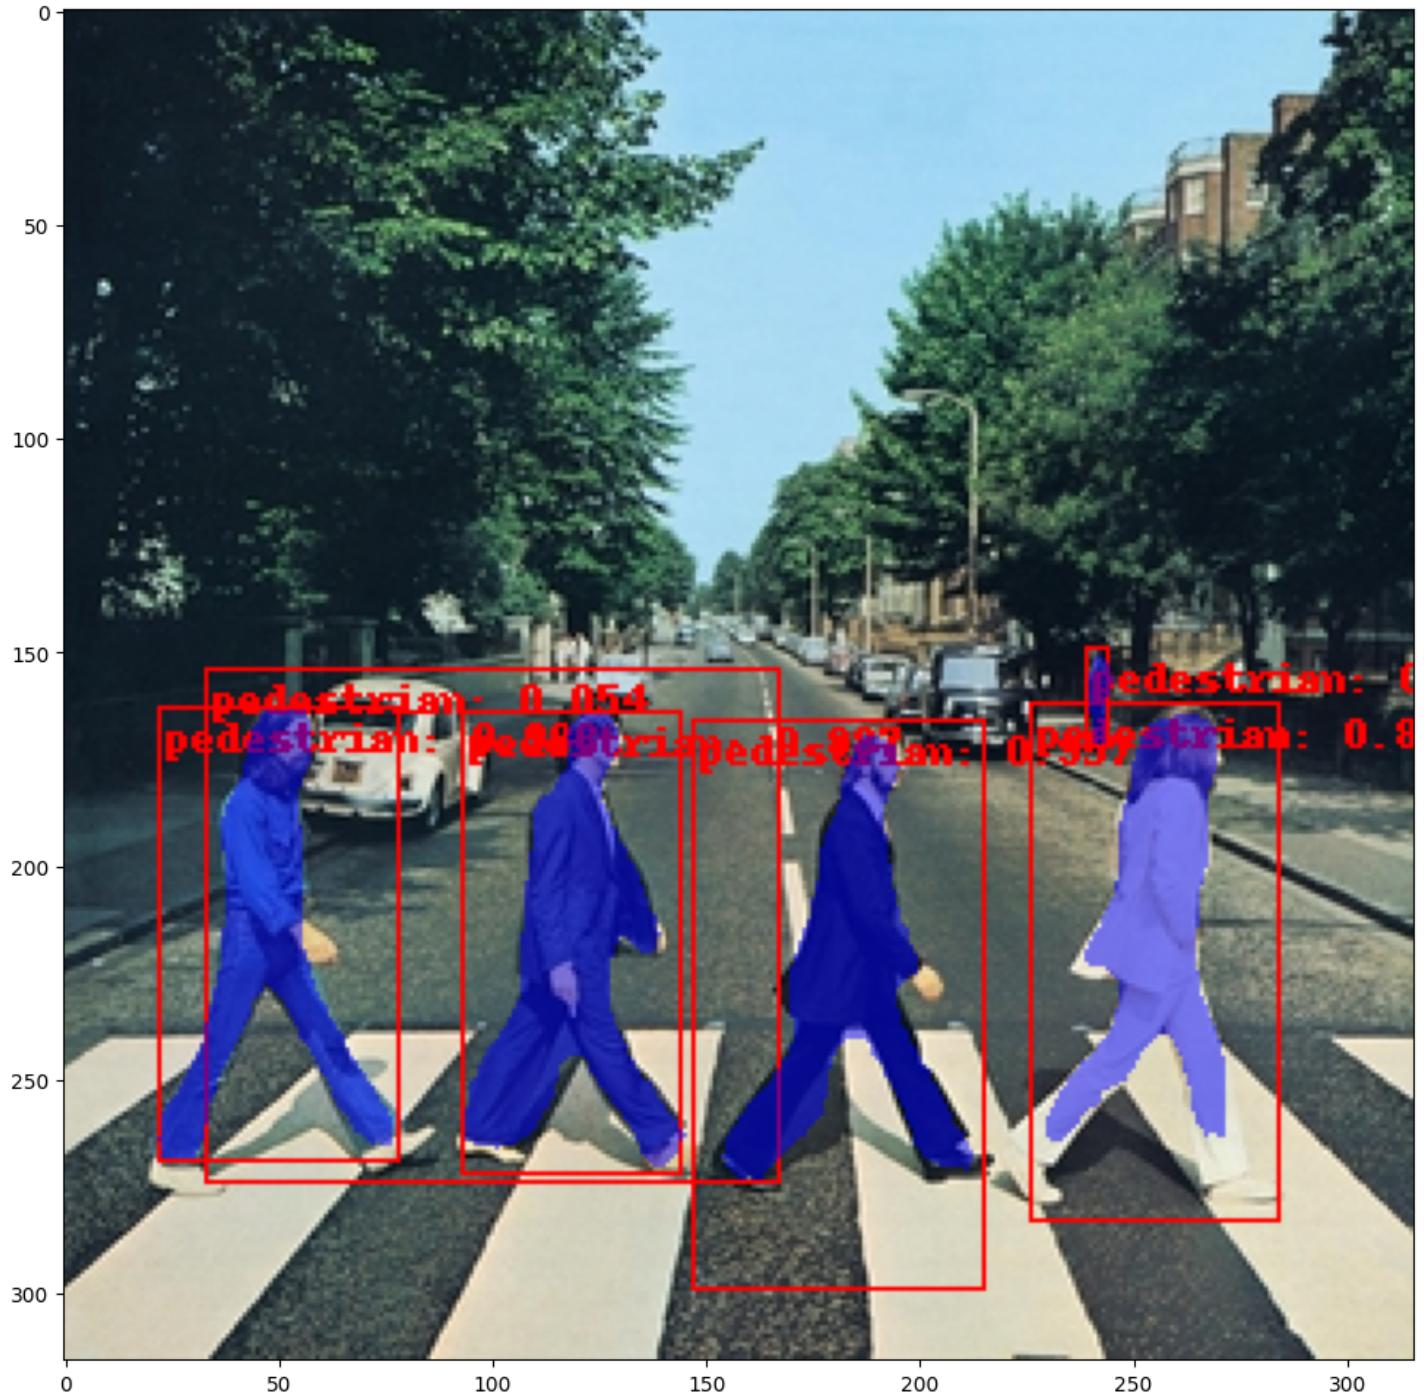
model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

image = (255.0 * (image - image.min()) / (image.max() - image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels, colors="red")

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5, colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))
```

```
<matplotlib.image.AxesImage at 0x79c685b5a6e0>
```



2 - Modifying the model to add a different backbone

```
import torchvision
from torchvision.models.detection import FasterRCNN, MaskRCNN
from torchvision.models.detection.rpn import AnchorGenerator

backbone = torchvision.models.mobilenet_v2(weights="DEFAULT").features
```

```
# ``FasterRCNN`` needs to know the number of
# output channels in a backbone. For mobilenet_v2, it's 1280
# so we need to add it here
backbone.out_channels = 1280

# let's make the RPN generate 5 x 3 anchors per spatial
# location, with 5 different sizes and 3 different aspect
# ratios. We have a Tuple[Tuple[int]] because each feature
# map could potentially have different sizes and
# aspect ratios
anchor_generator = AnchorGenerator(
    sizes=((32, 64, 128, 256, 512),),
    aspect_ratios=((0.5, 1.0, 2.0),)
)

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# ``OrderedDict[Tensor]``, and in ``featmap_names`` you can choose which
# feature maps to use.
roi_pooler = torchvision.ops.MultiScaleRoIAlign(
    featmap_names=['0'],
    output_size=7,
    sampling_ratio=2
)

# put the pieces together inside a Faster-RCNN model
model = FasterRCNN(
    backbone,
    num_classes=2,
    rpn_anchor_generator=anchor_generator,
    box_roi_pool=roi_pooler
)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
```

```

lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs, modify to 10
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

```

```

Epoch: [0]  [ 0/60]  eta: 0:00:23  lr: 0.000090  loss: 1.4881 (1.4881)  loss_classifier:
0.7213 (0.7213)  loss_box_reg: 0.0230 (0.0230)  loss_objectness: 0.7056 (0.7056)
loss_rpn_box_reg: 0.0382 (0.0382)  time: 0.3888  data: 0.0220  max mem: 13577
Epoch: [0]  [10/60]  eta: 0:00:17  lr: 0.000936  loss: 1.4354 (1.3951)  loss_classifier:
0.6727 (0.6331)  loss_box_reg: 0.0234 (0.0278)  loss_objectness: 0.6981 (0.6935)
loss_rpn_box_reg: 0.0385 (0.0408)  time: 0.3545  data: 0.0276  max mem: 13577
Epoch: [0]  [20/60]  eta: 0:00:14  lr: 0.001783  loss: 1.1219 (1.2181)  loss_classifier:
0.4221 (0.4668)  loss_box_reg: 0.0411 (0.0594)  loss_objectness: 0.6578 (0.6497)
loss_rpn_box_reg: 0.0406 (0.0422)  time: 0.3573  data: 0.0313  max mem: 13577
Epoch: [0]  [30/60]  eta: 0:00:10  lr: 0.002629  loss: 0.9818 (1.1438)  loss_classifier:
0.2722 (0.4175)  loss_box_reg: 0.1333 (0.0981)  loss_objectness: 0.5217 (0.5853)
loss_rpn_box_reg: 0.0479 (0.0429)  time: 0.3508  data: 0.0305  max mem: 13577
Epoch: [0]  [40/60]  eta: 0:00:06  lr: 0.003476  loss: 0.8954 (1.0571)  loss_classifier:
0.2722 (0.3807)  loss_box_reg: 0.1931 (0.1188)  loss_objectness: 0.3562 (0.5176)
loss_rpn_box_reg: 0.0340 (0.0399)  time: 0.3377  data: 0.0257  max mem: 13577
Epoch: [0]  [50/60]  eta: 0:00:03  lr: 0.004323  loss: 0.7315 (0.9852)  loss_classifier:
0.2473 (0.3516)  loss_box_reg: 0.1951 (0.1326)  loss_objectness: 0.2558 (0.4615)
loss_rpn_box_reg: 0.0297 (0.0394)  time: 0.3446  data: 0.0283  max mem: 13577
Epoch: [0]  [59/60]  eta: 0:00:00  lr: 0.005000  loss: 0.5636 (0.9258)  loss_classifier:
0.2049 (0.3316)  loss_box_reg: 0.1632 (0.1387)  loss_objectness: 0.2007 (0.4181)
loss_rpn_box_reg: 0.0254 (0.0374)  time: 0.3414  data: 0.0275  max mem: 13577
Epoch: [0] Total time: 0:00:20 (0.3468 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:04  model_time: 0.0805 (0.0805)  evaluator_time: 0.0014 (0.0014)
time: 0.0916  data: 0.0091  max mem: 13577
Test:  [49/50]  eta: 0:00:00  model_time: 0.0666 (0.0692)  evaluator_time: 0.0018 (0.0025)
time: 0.0818  data: 0.0127  max mem: 13577
Test: Total time: 0:00:04 (0.0854 s / it)

Averaged stats: model_time: 0.0666 (0.0692)  evaluator_time: 0.0018 (0.0025)
Accumulating evaluation results...

```

```

DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.061
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.188
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.010
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.001
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.111
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.060
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.281
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.338
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.022
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.368
Epoch: [1] [ 0/60] eta: 0:00:20 lr: 0.005000 loss: 0.7907 (0.7907) loss_classifier: 0.2730 (0.2730) loss_box_reg: 0.3139 (0.3139) loss_objectness: 0.1803 (0.1803) loss_rpn_box_reg: 0.0235 (0.0235) time: 0.3491 data: 0.0262 max mem: 13577
Epoch: [1] [10/60] eta: 0:00:17 lr: 0.005000 loss: 0.5971 (0.6338) loss_classifier: 0.2137 (0.2243) loss_box_reg: 0.2479 (0.2201) loss_objectness: 0.1608 (0.1604) loss_rpn_box_reg: 0.0235 (0.0290) time: 0.3586 data: 0.0342 max mem: 13577
Epoch: [1] [20/60] eta: 0:00:14 lr: 0.005000 loss: 0.5401 (0.5911) loss_classifier: 0.1893 (0.2035) loss_box_reg: 0.1889 (0.2079) loss_objectness: 0.1467 (0.1501) loss_rpn_box_reg: 0.0266 (0.0297) time: 0.3531 data: 0.0316 max mem: 13577
Epoch: [1] [30/60] eta: 0:00:10 lr: 0.005000 loss: 0.4755 (0.5633) loss_classifier: 0.1654 (0.1913) loss_box_reg: 0.1571 (0.1996) loss_objectness: 0.1304 (0.1437) loss_rpn_box_reg: 0.0254 (0.0287) time: 0.3436 data: 0.0276 max mem: 13577
Epoch: [1] [40/60] eta: 0:00:07 lr: 0.005000 loss: 0.4316 (0.5280) loss_classifier: 0.1458 (0.1799) loss_box_reg: 0.1333 (0.1885) loss_objectness: 0.1165 (0.1326) loss_rpn_box_reg: 0.0238 (0.0271) time: 0.3484 data: 0.0295 max mem: 13577
Epoch: [1] [50/60] eta: 0:00:03 lr: 0.005000 loss: 0.4568 (0.5262) loss_classifier: 0.1411 (0.1769) loss_box_reg: 0.1767 (0.1920) loss_objectness: 0.1101 (0.1293) loss_rpn_box_reg: 0.0267 (0.0279) time: 0.3489 data: 0.0286 max mem: 13577
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.4514 (0.5163) loss_classifier: 0.1388 (0.1722) loss_box_reg: 0.1767 (0.1905) loss_objectness: 0.1074 (0.1253) loss_rpn_box_reg: 0.0320 (0.0283) time: 0.3407 data: 0.0262 max mem: 13577
Epoch: [1] Total time: 0:00:20 (0.3473 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0709 (0.0709) evaluator_time: 0.0013 (0.0013) time: 0.0812 data: 0.0084 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0620 (0.0687) evaluator_time: 0.0014 (0.0023) time: 0.0781 data: 0.0125 max mem: 13577
Test: Total time: 0:00:04 (0.0864 s / it)
Averaged stats: model_time: 0.0620 (0.0687) evaluator_time: 0.0014 (0.0023)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.167
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.548

```

```

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.018
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.191
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.112
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.337
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.352
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.384
Epoch: [2] [ 0/60] eta: 0:00:22 lr: 0.005000 loss: 0.6456 (0.6456) loss_classifier:
0.1880 (0.1880) loss_box_reg: 0.2919 (0.2919) loss_objectness: 0.1256 (0.1256)
loss_rpn_box_reg: 0.0402 (0.0402) time: 0.3706 data: 0.0359 max mem: 13577
Epoch: [2] [10/60] eta: 0:00:17 lr: 0.005000 loss: 0.4584 (0.4519) loss_classifier:
0.1518 (0.1423) loss_box_reg: 0.2013 (0.1904) loss_objectness: 0.0873 (0.0925)
loss_rpn_box_reg: 0.0289 (0.0267) time: 0.3464 data: 0.0262 max mem: 13577
Epoch: [2] [20/60] eta: 0:00:13 lr: 0.005000 loss: 0.4640 (0.4822) loss_classifier:
0.1518 (0.1485) loss_box_reg: 0.2013 (0.2098) loss_objectness: 0.0873 (0.0966)
loss_rpn_box_reg: 0.0254 (0.0272) time: 0.3455 data: 0.0264 max mem: 13577
Epoch: [2] [30/60] eta: 0:00:10 lr: 0.005000 loss: 0.4017 (0.4432) loss_classifier:
0.1267 (0.1367) loss_box_reg: 0.1885 (0.1936) loss_objectness: 0.0834 (0.0878)
loss_rpn_box_reg: 0.0203 (0.0252) time: 0.3533 data: 0.0308 max mem: 13577
Epoch: [2] [40/60] eta: 0:00:07 lr: 0.005000 loss: 0.3251 (0.4346) loss_classifier:
0.1002 (0.1323) loss_box_reg: 0.1437 (0.1879) loss_objectness: 0.0703 (0.0867)
loss_rpn_box_reg: 0.0215 (0.0278) time: 0.3556 data: 0.0334 max mem: 13577
Epoch: [2] [50/60] eta: 0:00:03 lr: 0.005000 loss: 0.3320 (0.4233) loss_classifier:
0.0987 (0.1296) loss_box_reg: 0.1368 (0.1840) loss_objectness: 0.0687 (0.0826)
loss_rpn_box_reg: 0.0266 (0.0272) time: 0.3495 data: 0.0288 max mem: 13577
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2990 (0.4011) loss_classifier:
0.0873 (0.1230) loss_box_reg: 0.1261 (0.1743) loss_objectness: 0.0560 (0.0782)
loss_rpn_box_reg: 0.0202 (0.0256) time: 0.3484 data: 0.0252 max mem: 13577
Epoch: [2] Total time: 0:00:21 (0.3503 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0782 (0.0782) evaluator_time: 0.0013 (0.0013)
time: 0.0885 data: 0.0084 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0665 (0.0678) evaluator_time: 0.0015 (0.0019)
time: 0.0814 data: 0.0127 max mem: 13577
Test: Total time: 0:00:04 (0.0835 s / it)
Averaged stats: model_time: 0.0665 (0.0678) evaluator_time: 0.0015 (0.0019)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.172
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.555
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.032
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.190

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.108
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.364
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.424
Epoch: [3] [ 0/60] eta: 0:00:22 lr: 0.000500 loss: 0.6482 (0.6482) loss_classifier:
0.1973 (0.1973) loss_box_reg: 0.3128 (0.3128) loss_objectness: 0.0833 (0.0833)
loss_rpn_box_reg: 0.0548 (0.0548) time: 0.3675 data: 0.0290 max mem: 13577
Epoch: [3] [10/60] eta: 0:00:18 lr: 0.000500 loss: 0.4065 (0.4286) loss_classifier:
0.1095 (0.1292) loss_box_reg: 0.2072 (0.2068) loss_objectness: 0.0707 (0.0688)
loss_rpn_box_reg: 0.0205 (0.0238) time: 0.3614 data: 0.0307 max mem: 13577
Epoch: [3] [20/60] eta: 0:00:14 lr: 0.000500 loss: 0.3403 (0.3605) loss_classifier:
0.0980 (0.1110) loss_box_reg: 0.1423 (0.1704) loss_objectness: 0.0562 (0.0607)
loss_rpn_box_reg: 0.0155 (0.0183) time: 0.3550 data: 0.0291 max mem: 13577
Epoch: [3] [30/60] eta: 0:00:10 lr: 0.000500 loss: 0.3084 (0.3541) loss_classifier:
0.0932 (0.1080) loss_box_reg: 0.1388 (0.1659) loss_objectness: 0.0537 (0.0604)
loss_rpn_box_reg: 0.0138 (0.0198) time: 0.3498 data: 0.0274 max mem: 13577
Epoch: [3] [40/60] eta: 0:00:07 lr: 0.000500 loss: 0.2729 (0.3315) loss_classifier:
0.0804 (0.1009) loss_box_reg: 0.1281 (0.1543) loss_objectness: 0.0518 (0.0577)
loss_rpn_box_reg: 0.0186 (0.0186) time: 0.3528 data: 0.0281 max mem: 13577
Epoch: [3] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.2782 (0.3425) loss_classifier:
0.0804 (0.1032) loss_box_reg: 0.1396 (0.1606) loss_objectness: 0.0470 (0.0583)
loss_rpn_box_reg: 0.0164 (0.0204) time: 0.3660 data: 0.0325 max mem: 13577
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.3178 (0.3415) loss_classifier:
0.1016 (0.1030) loss_box_reg: 0.1470 (0.1588) loss_objectness: 0.0548 (0.0590)
loss_rpn_box_reg: 0.0212 (0.0206) time: 0.3662 data: 0.0307 max mem: 13577
Epoch: [3] Total time: 0:00:21 (0.3575 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0784 (0.0784) evaluator_time: 0.0013 (0.0013)
time: 0.0889 data: 0.0086 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0651 (0.0665) evaluator_time: 0.0015 (0.0020)
time: 0.0813 data: 0.0131 max mem: 13577
Test: Total time: 0:00:04 (0.0825 s / it)
Averaged stats: model_time: 0.0651 (0.0665) evaluator_time: 0.0015 (0.0020)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.247
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.673
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.083
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.001
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.276
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.168
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.390
Average Recall      (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.408
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.022
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.445
Epoch: [4] [ 0/60] eta: 0:00:19 lr: 0.000500 loss: 0.2327 (0.2327) loss_classifier:
0.0778 (0.0778) loss_box_reg: 0.1046 (0.1046) loss_objectness: 0.0390 (0.0390)
loss_rpn_box_reg: 0.0114 (0.0114) time: 0.3206 data: 0.0183 max mem: 13577
Epoch: [4] [10/60] eta: 0:00:18 lr: 0.000500 loss: 0.3444 (0.3246) loss_classifier:
0.0968 (0.0970) loss_box_reg: 0.1720 (0.1603) loss_objectness: 0.0466 (0.0502)
loss_rpn_box_reg: 0.0146 (0.0172) time: 0.3659 data: 0.0314 max mem: 13577
Epoch: [4] [20/60] eta: 0:00:14 lr: 0.000500 loss: 0.2942 (0.3074) loss_classifier:
0.0941 (0.0938) loss_box_reg: 0.1415 (0.1474) loss_objectness: 0.0446 (0.0483)
loss_rpn_box_reg: 0.0148 (0.0178) time: 0.3607 data: 0.0302 max mem: 13577
Epoch: [4] [30/60] eta: 0:00:10 lr: 0.000500 loss: 0.2816 (0.3297) loss_classifier:
0.0858 (0.0991) loss_box_reg: 0.1317 (0.1595) loss_objectness: 0.0447 (0.0517)
loss_rpn_box_reg: 0.0178 (0.0193) time: 0.3529 data: 0.0272 max mem: 13577
Epoch: [4] [40/60] eta: 0:00:07 lr: 0.000500 loss: 0.3514 (0.3424) loss_classifier:
0.1076 (0.1034) loss_box_reg: 0.1787 (0.1645) loss_objectness: 0.0569 (0.0547)
loss_rpn_box_reg: 0.0214 (0.0198) time: 0.3623 data: 0.0299 max mem: 13577
Epoch: [4] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.3037 (0.3354) loss_classifier:
0.1046 (0.1019) loss_box_reg: 0.1417 (0.1597) loss_objectness: 0.0537 (0.0545)
loss_rpn_box_reg: 0.0183 (0.0193) time: 0.3596 data: 0.0287 max mem: 13577
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.2808 (0.3273) loss_classifier:
0.0823 (0.0992) loss_box_reg: 0.1188 (0.1551) loss_objectness: 0.0519 (0.0534)
loss_rpn_box_reg: 0.0173 (0.0196) time: 0.3518 data: 0.0253 max mem: 13577
Epoch: [4] Total time: 0:00:21 (0.3577 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0780 (0.0780) evaluator_time: 0.0012 (0.0012)
time: 0.0882 data: 0.0084 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0647 (0.0713) evaluator_time: 0.0018 (0.0028)
time: 0.0810 data: 0.0130 max mem: 13577
Test: Total time: 0:00:04 (0.0898 s / it)
Averaged stats: model_time: 0.0647 (0.0713) evaluator_time: 0.0018 (0.0028)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.238
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.609
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.132
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.001
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.265
Average Recall   (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.157
Average Recall   (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.398
Average Recall   (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.425
Average Recall   (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall   (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.011
Average Recall   (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.463
Epoch: [5] [ 0/60] eta: 0:00:20 lr: 0.000500 loss: 0.3760 (0.3760) loss_classifier:
0.1191 (0.1191) loss_box_reg: 0.1695 (0.1695) loss_objectness: 0.0636 (0.0636)

```

```

loss_rpn_box_reg: 0.0238 (0.0238) time: 0.3363 data: 0.0218 max mem: 13577
Epoch: [5] [10/60] eta: 0:00:17 lr: 0.000500 loss: 0.3568 (0.3450) loss_classifier:
0.0978 (0.1029) loss_box_reg: 0.1695 (0.1675) loss_objectness: 0.0550 (0.0541)
loss_rpn_box_reg: 0.0179 (0.0204) time: 0.3514 data: 0.0271 max mem: 13577
Epoch: [5] [20/60] eta: 0:00:13 lr: 0.000500 loss: 0.3011 (0.3082) loss_classifier:
0.0862 (0.0902) loss_box_reg: 0.1370 (0.1482) loss_objectness: 0.0473 (0.0517)
loss_rpn_box_reg: 0.0152 (0.0181) time: 0.3490 data: 0.0265 max mem: 13577
Epoch: [5] [30/60] eta: 0:00:10 lr: 0.000500 loss: 0.3011 (0.3117) loss_classifier:
0.0825 (0.0910) loss_box_reg: 0.1342 (0.1498) loss_objectness: 0.0473 (0.0520)
loss_rpn_box_reg: 0.0171 (0.0189) time: 0.3564 data: 0.0286 max mem: 13577
Epoch: [5] [40/60] eta: 0:00:07 lr: 0.000500 loss: 0.3152 (0.3197) loss_classifier:
0.0937 (0.0942) loss_box_reg: 0.1516 (0.1533) loss_objectness: 0.0534 (0.0531)
loss_rpn_box_reg: 0.0205 (0.0192) time: 0.3633 data: 0.0287 max mem: 13577
Epoch: [5] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.2916 (0.3270) loss_classifier:
0.0986 (0.0979) loss_box_reg: 0.1373 (0.1551) loss_objectness: 0.0516 (0.0545)
loss_rpn_box_reg: 0.0163 (0.0194) time: 0.3611 data: 0.0273 max mem: 13577
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.3065 (0.3231) loss_classifier:
0.0896 (0.0967) loss_box_reg: 0.1374 (0.1539) loss_objectness: 0.0418 (0.0527)
loss_rpn_box_reg: 0.0167 (0.0198) time: 0.3660 data: 0.0308 max mem: 13577
Epoch: [5] Total time: 0:00:21 (0.3597 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0827 (0.0827) evaluator_time: 0.0013 (0.0013)
time: 0.0933 data: 0.0087 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0639 (0.0663) evaluator_time: 0.0016 (0.0020)
time: 0.0792 data: 0.0125 max mem: 13577
Test: Total time: 0:00:04 (0.0818 s / it)
Averaged stats: model_time: 0.0639 (0.0663) evaluator_time: 0.0016 (0.0020)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.269
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.637
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.137
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.001
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.301
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.160
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.419
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.442
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.033
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.480
Epoch: [6] [ 0/60] eta: 0:00:19 lr: 0.000050 loss: 0.1891 (0.1891) loss_classifier:
0.0685 (0.0685) loss_box_reg: 0.0786 (0.0786) loss_objectness: 0.0355 (0.0355)
loss_rpn_box_reg: 0.0065 (0.0065) time: 0.3305 data: 0.0256 max mem: 13577

Epoch: [6] [10/60] eta: 0:00:17 lr: 0.000050 loss: 0.3023 (0.3003) loss_classifier:
0.0823 (0.0958) loss_box_reg: 0.1401 (0.1395) loss_objectness: 0.0473 (0.0460)

```

```

loss_rpn_box_reg: 0.0194 (0.0189) time: 0.3567 data: 0.0293 max mem: 13577
Epoch: [6] [20/60] eta: 0:00:14 lr: 0.000050 loss: 0.2657 (0.3032) loss_classifier:
0.0823 (0.0958) loss_box_reg: 0.1380 (0.1448) loss_objectness: 0.0437 (0.0449)
loss_rpn_box_reg: 0.0178 (0.0177) time: 0.3575 data: 0.0290 max mem: 13577
Epoch: [6] [30/60] eta: 0:00:10 lr: 0.000050 loss: 0.2603 (0.3012) loss_classifier:
0.0834 (0.0944) loss_box_reg: 0.1290 (0.1412) loss_objectness: 0.0457 (0.0476)
loss_rpn_box_reg: 0.0160 (0.0181) time: 0.3573 data: 0.0260 max mem: 13577
Epoch: [6] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.2958 (0.3150) loss_classifier:
0.0929 (0.0985) loss_box_reg: 0.1340 (0.1493) loss_objectness: 0.0474 (0.0479)
loss_rpn_box_reg: 0.0171 (0.0193) time: 0.3646 data: 0.0266 max mem: 13577
Epoch: [6] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.3133 (0.3261) loss_classifier:
0.0961 (0.0998) loss_box_reg: 0.1451 (0.1554) loss_objectness: 0.0474 (0.0511)
loss_rpn_box_reg: 0.0216 (0.0198) time: 0.3688 data: 0.0321 max mem: 13577
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.2958 (0.3167) loss_classifier:
0.0879 (0.0968) loss_box_reg: 0.1273 (0.1506) loss_objectness: 0.0440 (0.0496)
loss_rpn_box_reg: 0.0201 (0.0197) time: 0.3575 data: 0.0289 max mem: 13577
Epoch: [6] Total time: 0:00:21 (0.3597 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0778 (0.0778) evaluator_time: 0.0014 (0.0014)
time: 0.0896 data: 0.0098 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0676 (0.0673) evaluator_time: 0.0019 (0.0022)
time: 0.0865 data: 0.0146 max mem: 13577
Test: Total time: 0:00:04 (0.0840 s / it)
Averaged stats: model_time: 0.0676 (0.0673) evaluator_time: 0.0019 (0.0022)
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.256
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.632
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.135
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.016
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.289
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.164
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.414
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.454
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.044
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.492
Epoch: [7] [ 0/60] eta: 0:00:22 lr: 0.000050 loss: 0.3784 (0.3784) loss_classifier:
0.1157 (0.1157) loss_box_reg: 0.1815 (0.1815) loss_objectness: 0.0602 (0.0602)
loss_rpn_box_reg: 0.0211 (0.0211) time: 0.3757 data: 0.0321 max mem: 13577
Epoch: [7] [10/60] eta: 0:00:18 lr: 0.000050 loss: 0.3713 (0.3509) loss_classifier:
0.1037 (0.1044) loss_box_reg: 0.1815 (0.1706) loss_objectness: 0.0561 (0.0538)
loss_rpn_box_reg: 0.0211 (0.0221) time: 0.3666 data: 0.0316 max mem: 13577

Epoch: [7] [20/60] eta: 0:00:14 lr: 0.000050 loss: 0.2973 (0.3345) loss_classifier:
0.0902 (0.1013) loss_box_reg: 0.1461 (0.1609) loss_objectness: 0.0478 (0.0516)

```

```

loss_rpn_box_reg: 0.0181 (0.0207) time: 0.3618 data: 0.0293 max mem: 13577
Epoch: [7] [30/60] eta: 0:00:10 lr: 0.000050 loss: 0.3101 (0.3373) loss_classifier:
0.0902 (0.1041) loss_box_reg: 0.1538 (0.1602) loss_objectness: 0.0503 (0.0521)
loss_rpn_box_reg: 0.0181 (0.0209) time: 0.3556 data: 0.0256 max mem: 13577
Epoch: [7] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.3186 (0.3300) loss_classifier:
0.0974 (0.1008) loss_box_reg: 0.1454 (0.1577) loss_objectness: 0.0514 (0.0517)
loss_rpn_box_reg: 0.0162 (0.0198) time: 0.3558 data: 0.0286 max mem: 13577
Epoch: [7] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.2648 (0.3152) loss_classifier:
0.0812 (0.0961) loss_box_reg: 0.1153 (0.1500) loss_objectness: 0.0435 (0.0504)
loss_rpn_box_reg: 0.0112 (0.0187) time: 0.3572 data: 0.0298 max mem: 13577
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.2703 (0.3175) loss_classifier:
0.0835 (0.0961) loss_box_reg: 0.1105 (0.1514) loss_objectness: 0.0430 (0.0507)
loss_rpn_box_reg: 0.0127 (0.0194) time: 0.3554 data: 0.0271 max mem: 13577
Epoch: [7] Total time: 0:00:21 (0.3586 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0841 (0.0841) evaluator_time: 0.0022 (0.0022)
time: 0.0987 data: 0.0117 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0614 (0.0672) evaluator_time: 0.0016 (0.0023)
time: 0.0778 data: 0.0126 max mem: 13577
Test: Total time: 0:00:04 (0.0842 s / it)
Averaged stats: model_time: 0.0614 (0.0672) evaluator_time: 0.0016 (0.0023)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.255
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.658
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.107
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.001
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.285
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.177
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.427
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.452
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.033
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.492
Epoch: [8] [ 0/60] eta: 0:00:21 lr: 0.000050 loss: 0.2255 (0.2255) loss_classifier:
0.0732 (0.0732) loss_box_reg: 0.1081 (0.1081) loss_objectness: 0.0320 (0.0320)
loss_rpn_box_reg: 0.0122 (0.0122) time: 0.3532 data: 0.0234 max mem: 13577
Epoch: [8] [10/60] eta: 0:00:17 lr: 0.000050 loss: 0.2585 (0.2776) loss_classifier:
0.0870 (0.0859) loss_box_reg: 0.1270 (0.1322) loss_objectness: 0.0390 (0.0416)
loss_rpn_box_reg: 0.0140 (0.0179) time: 0.3519 data: 0.0271 max mem: 13577
Epoch: [8] [20/60] eta: 0:00:14 lr: 0.000050 loss: 0.2585 (0.2887) loss_classifier:
0.0882 (0.0894) loss_box_reg: 0.1270 (0.1377) loss_objectness: 0.0411 (0.0443)
loss_rpn_box_reg: 0.0140 (0.0174) time: 0.3521 data: 0.0271 max mem: 13577

Epoch: [8] [30/60] eta: 0:00:10 lr: 0.000050 loss: 0.2462 (0.3003) loss_classifier:
0.0859 (0.0912) loss_box_reg: 0.1124 (0.1417) loss_objectness: 0.0479 (0.0482)

```

```

loss_rpn_box_reg: 0.0128 (0.0191) time: 0.3544 data: 0.0212 max mem: 13577
Epoch: [8] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.3002 (0.3212) loss_classifier:
0.0921 (0.0971) loss_box_reg: 0.1475 (0.1531) loss_objectness: 0.0475 (0.0509)
loss_rpn_box_reg: 0.0183 (0.0202) time: 0.3622 data: 0.0288 max mem: 13577
Epoch: [8] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.3099 (0.3193) loss_classifier:
0.0986 (0.0962) loss_box_reg: 0.1511 (0.1521) loss_objectness: 0.0475 (0.0508)
loss_rpn_box_reg: 0.0190 (0.0201) time: 0.3685 data: 0.0298 max mem: 13577
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.3078 (0.3136) loss_classifier:
0.0859 (0.0950) loss_box_reg: 0.1388 (0.1489) loss_objectness: 0.0431 (0.0502)
loss_rpn_box_reg: 0.0170 (0.0195) time: 0.3680 data: 0.0306 max mem: 13577
Epoch: [8] Total time: 0:00:21 (0.3604 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0798 (0.0798) evaluator_time: 0.0012 (0.0012)
time: 0.0901 data: 0.0086 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0641 (0.0668) evaluator_time: 0.0016 (0.0020)
time: 0.0796 data: 0.0126 max mem: 13577
Test: Total time: 0:00:04 (0.0825 s / it)
Averaged stats: model_time: 0.0641 (0.0668) evaluator_time: 0.0016 (0.0020)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.260
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.668
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.110
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.012
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.287
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.190
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.412
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.427
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.022
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.465
Epoch: [9] [ 0/60] eta: 0:00:22 lr: 0.000005 loss: 0.3433 (0.3433) loss_classifier:
0.1244 (0.1244) loss_box_reg: 0.1495 (0.1495) loss_objectness: 0.0454 (0.0454)
loss_rpn_box_reg: 0.0239 (0.0239) time: 0.3696 data: 0.0260 max mem: 13577
Epoch: [9] [10/60] eta: 0:00:18 lr: 0.000005 loss: 0.3433 (0.3450) loss_classifier:
0.1176 (0.1109) loss_box_reg: 0.1586 (0.1660) loss_objectness: 0.0454 (0.0504)
loss_rpn_box_reg: 0.0179 (0.0178) time: 0.3737 data: 0.0329 max mem: 13577
Epoch: [9] [20/60] eta: 0:00:14 lr: 0.000005 loss: 0.2860 (0.3264) loss_classifier:
0.0818 (0.1020) loss_box_reg: 0.1315 (0.1515) loss_objectness: 0.0453 (0.0546)
loss_rpn_box_reg: 0.0165 (0.0183) time: 0.3614 data: 0.0295 max mem: 13577
Epoch: [9] [30/60] eta: 0:00:10 lr: 0.000005 loss: 0.2700 (0.3164) loss_classifier:
0.0832 (0.0982) loss_box_reg: 0.1310 (0.1485) loss_objectness: 0.0417 (0.0515)
loss_rpn_box_reg: 0.0178 (0.0181) time: 0.3585 data: 0.0274 max mem: 13577

Epoch: [9] [40/60] eta: 0:00:07 lr: 0.000005 loss: 0.2760 (0.3095) loss_classifier:
0.0834 (0.0960) loss_box_reg: 0.1424 (0.1460) loss_objectness: 0.0359 (0.0495)
loss_rpn_box_reg: 0.0179 (0.0179) time: 0.3593 data: 0.0271 max mem: 13577

```

```

loss_rpn_box_reg: 0.0110 (0.0117) time: 0.3505 data: 0.0281 max mem: 13577
Epoch: [9] [50/60] eta: 0:00:03 lr: 0.000005 loss: 0.2891 (0.3193) loss_classifier:
0.0834 (0.0970) loss_box_reg: 0.1424 (0.1502) loss_objectness: 0.0457 (0.0518)
loss_rpn_box_reg: 0.0194 (0.0203) time: 0.3616 data: 0.0283 max mem: 13577
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.2806 (0.3094) loss_classifier:
0.0703 (0.0939) loss_box_reg: 0.1424 (0.1463) loss_objectness: 0.0421 (0.0499)
loss_rpn_box_reg: 0.0175 (0.0193) time: 0.3663 data: 0.0288 max mem: 13577
Epoch: [9] Total time: 0:00:21 (0.3625 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0782 (0.0782) evaluator_time: 0.0013 (0.0013)
time: 0.0885 data: 0.0084 max mem: 13577
Test: [49/50] eta: 0:00:00 model_time: 0.0707 (0.0696) evaluator_time: 0.0023 (0.0025)
time: 0.0906 data: 0.0161 max mem: 13577
Test: Total time: 0:00:04 (0.0876 s / it)
Averaged stats: model_time: 0.0707 (0.0696) evaluator_time: 0.0023 (0.0025)
Accumulating evaluation results...
DONE (t=0.03s).

IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.257
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.688
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.133
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.013
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.287
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.164
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.422
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.441
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.044
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.478

```

```

import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes, draw_segmentation_masks

image = read_image("Beatles_-_Abbey_Road.jpg")
eval_transform = get_transform(train=False)

model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

```

```
preds = predictions[0]

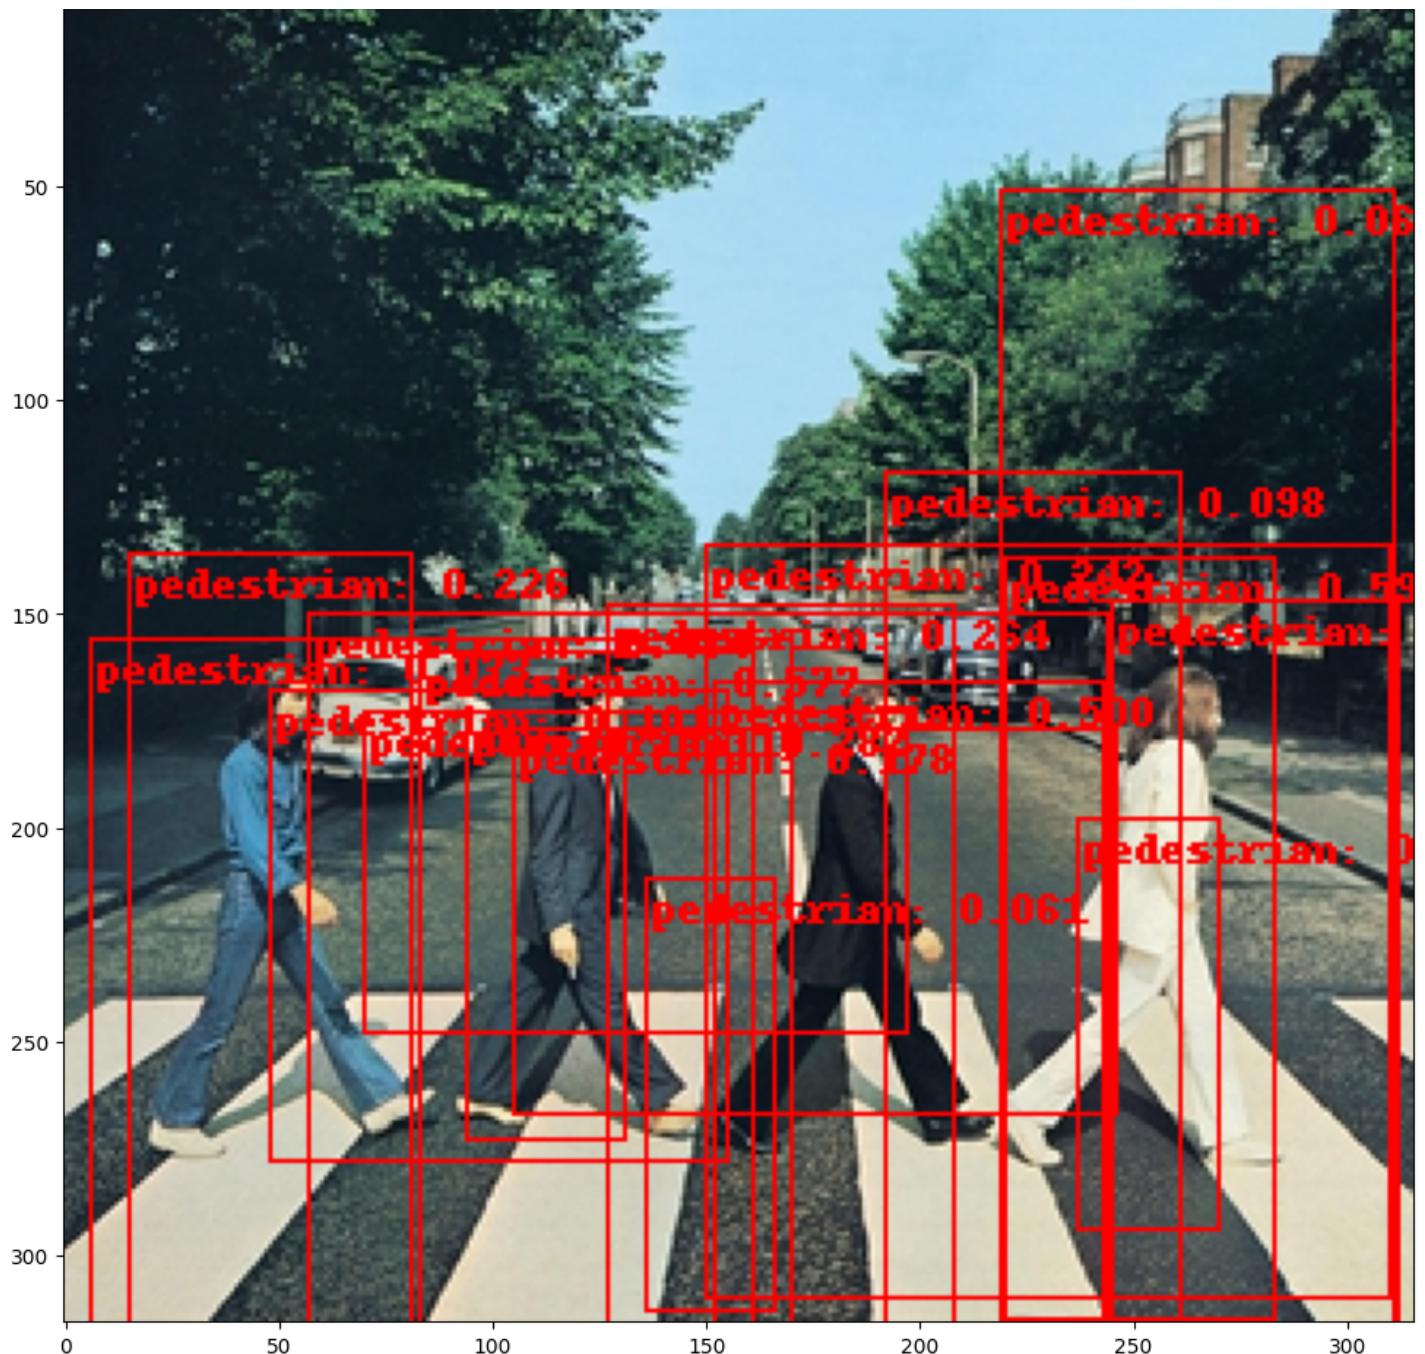
image = (255.0 * (image - image.min()) / (image.max() - image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels, colors="red")

# masks = (pred["masks"] > 0.7).squeeze(1)
# output_image = draw_segmentation_masks(output_image, masks, alpha=0.5, colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))
```

```
<matplotlib.image.AxesImage at 0x79c67ea2ed70>
```





(b). On the training data, option 1 perform much better than option 2.

(c). On the testing picture, option 1 also perform much better than option 2. Option 1 is a MaskRCNN model, it not only detects objects, gives their bounding boxes, but also provides a pixel-wise mask for each detected object. This mask can differentiate pixels belonging to the object from the background.